The Bitcoin Network

Implementation

The only source you really need:

https://en.bitcoin.it/wiki/Protocol documentation

Simple explanation of the network protocol/messages

But let's go step by step

Bitcoin nodes:

- Run software implementing the logic of Bitcoin (transactions, blockchain, proof of work, addresses of other nodes)
- They create a p2p network for transmitting messages (gossip protocol)
- Messages can also be interchanged outside of this network!!!

Objective of this class:

 Implement (part of) the communication bewteen the nodes (be it full, SPV, or users that do not run a node)

What a message looks like?

- f9beb4d9 network magic (always 0xf9beb4d9 for mainnet)
- 76657273696f6e0000000000 command, 12 bytes, human-readable
- 65000000 payload length, 4 bytes, little-endian
- 5f1a69d2 payload checksum, first 4 bytes of hash256 of the payload
- 7211...01 payload

- Magic (4 bytes):
- Where message starts
- If the connection fails
- Identifies the network
- 0x0b110907 for testnet

-332fcf0

an looks like?

- f9beb4d9 network magic (always 0xf9beb4d9 for mainnet)
- 76657273696f6e0000000000 command, 12 bytes, human-readable
- 65000000 payload length, 4 bytes, little-endian
- 5f1a69d2 payload checksum, first 4 bytes of hash256 of the payload
- 7211...01 payload

- Command (12 bytes):
- e.g. 'block', 'header',...
 - ASCII string
- https://en.bitcoin.it/wiki/Protocol_d ocumentation

a looks like?

- f9beb4d9 network magic (alway 0xf9beb4d9 for mainnet)
- 76657273696f6e0000000000 command, 12 bytes, human-readable
- 65000000 payload length, 4 bytes, little-endian
- 5f1a69d2 payload checksum, first 4 bytes of hash256 of the payload
- 7211...01 payload



- f9beb4d9 network magic (a) s 0xf9beb4d9 for mainnet)
- 76657273696f6e0000000000 command, 12 bytes, human-readable
- 65000000 payload length, 4 bytes, little-endian
- 5f1a69d2 payload checksum, first 4 bytes of hash256 of the payload
- 7211...01 payload

looks like?

- f9beb4d9 network magic (0xf9beb4d9 for mainnet)
- 76657273696f6e0000000000 command, 12 bytes, human-readable
- 65000000 payload length, 4 bytes, little-endian
- 5f1a69d2 payload checksum, first 4 bytes of hash256 of the payload
- 7211...01 payload

The meat!

##

- f9beb4d9 network magic (alway 9beb4d9 for mainnet)
- 65000000 payload length bytes, little-endian
- 5f1a69d2 payload chee sum, first 4 bytes of hash256 of the payload
- 7211...01 payload

Implementation - class that stores the message

```
class NetworkEnvelope:
    def __init__(self, command, payload, testnet=False):
        self.command = command
        self.payload = payload
        if testnet:
            self.magic = TESTNET NETWORK MAGIC
        else:
            self.magic = NETWORK MAGIC
```

Implementation - class that stores the message

```
clase Notwork Envolume.
          oclassmethod
         def parse(cls, s, testnet=False):
                                                                                                            lse):
             magic = s.read(4)
             if magic == b'':
                 raise RuntimeError('Connection reset!')
             if testnet:
                 expected magic = TESTNET NETWORK MAGIC
                 expected magic = NETWORK MAGIC
             if magic != expected magic:
                 raise RuntimeError('magic is not right {} vs {}'.format(magic.hex(), expected magic.hex()))
              command = s.read(12)
              command = command.strip(b' \times 00')
              payload length = little endian to int(s.read(4))
              checksum = s.read(4)
              payload = s.read(payload_length)
              calculated_checksum = hash256(payload)[:4]
              if calculated checksum != checksum:
                 raise RuntimeError('checksum does not match')
             return cls(command, payload, testnet=testnet)
```

Implementation - class that stores the message

class NetworkEnvelope:

```
def serialize(self):
    '''Returns the byte serialization of the entire network message'''
    # add the network magic
    result = self.magic
    # command 12 bytes
    # fill with 0's
    result += self.command + b'\x00' * (12 - len(self.command))
    # payload length 4 bytes, little endian
    result += int to little endian(len(self.payload), 4)
    # checksum 4 bytes, first four of hash256 of payload
    result += hash256(self.payload)[:4]
    # payload
    result += self.payload
    return result
```

Some important commands

In this class we will implement:

- Version
- Verack
- GetHeaders
- Headers

In our implementation:

- To send a message: we need serialize
- To process a received message: we need parse

Version

```
2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f0000000001
7f110100 - Protocol version, 4 bytes, little-endian, 70015
000000000000000 - Network services of sender, 8 bytes, little-endian
ad17835b00000000 - Timestamp, 8 bytes, little-endian
000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
       0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
0000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
       0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
       /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

Version serves to establish a connection between two nodes!

Version

f6a8d7a440ec27a11b

JOOO01

```
000000000000000 - Network services of sender, 8 bytes, little-endian
ad17835b00000000 - Timestamp, 8 bytes, little-endian
000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

7f110100 - Protocol version, 4 bytes, little-endian, 70015

What type of message can we interchange?

Version

we interchange? 0000000000000ffff0 f6a8d7a440ec27a11b

```
7f110100 - Protocol version, 4 bytes, little-endian, 70015
000000000000000 - Network services of sender, 8 bytes, little-endian
ad17835b00000000 - Timestamp, 8 bytes, little-endian
000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

What protocol does the node use (segWit, Litecoin, BCH)?

Version

J00001

7f110100 - Protocol version, 4 bytes, little-endian, 70015 000000000000000 - Network services of sender, 8 bytes, little-endian ad17835b00000000 - Timestamp, 8 bytes, little-endian 000000000000000 - Network services of receiver, 8 bytes, little-endian 00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4 0.0.0.0 8d20 - Network port of receiver, 2 bytes, 8333 000000000000000 - Network services of sender, 8 bytes, little-endian 00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4 0.0.0.0 8d20 - Network port of sender, 2 bytes, 8333 f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses 1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent /programmingblockchain:0.1/ 00000000 - Height, 0 01 - Optional flag for relay, based on BIP37

What type of stuff do I need for this connection

Version

 need for this connection 000000000000000ffff0 f6a8d7a440ec27a11b

```
ad17835b00000000 - Timestamp, 8 bytes, little-endian
000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

7f110100 - Protocol versi, 4 bytes, little-endian, 70015

000000000000000 - Network services of sender, 8 bytes, little-endian

000000000000ffff0

6a8d7a440ec27a11b

JU0001

What is the time?
Serves to accept/reject
blocks (and for other things)

Version

00000000 - Height, 0

01 - Optional flag for relay, based on BIP37

```
7f110100 - Protocol version ytes, little-endian, 70015
000000000000000 - Networkservices of sender, 8 bytes, little-endian
ad17835b00000000 - Timestamp, 8 bytes, little-endian
000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
```

000000000000ffff0

Address to which we send the messages

Version

```
00000008d200000000000000000
                                                                6a8d7a440ec27a11b
2f70726f6772616d6d696e67626c6f6
                                                          J00001
7f110100 - Protocol version, 4///, little-endian, 70015
000000000000000 - Network secs of sender, 8 bytes, little-endian
ad17835b00000000 - Timestam 8 bytes, little-endian
0000000000000000 - Network services of receiver, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
0000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

000000000000ffff0

Address of the node

Cle-endian, 70015

Zender, 8 bytes, little-endian

Version

```
sending the message
00000008d200000000000000000
                                                   6a8d7a440ec27a11b
2f70726f6772616d6d696e67626c6f6
                                              00001
```

7f110100 - Protocol version, 4 byte

00000000000000000 - Network service

```
little-endian
ad17835b000000000 - Timestamp, 8 b
                                    of receiver, 8 bytes, little-endian
00000000000000000 - Network servi
                                   Network address of receiver, 16 bytes, IPv4
000000000000000000000ffff0000000
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
000000000000000 - Network services of sender, 8 bytes, little-endian
00000000000000000000ffff00000000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

000000000000ffff0

To detect if we are connecting to ourselves

Version

```
6a8d7a440ec27a11b
00000008d200000000000000000
2f70726f6772616d6d696e67626c6f6
                                                          J00001
7f110100 - Protocol version, 4 byte
                                         Zle-endian, 70015
00000000000000000 - Network service
                                       Mender, 8 bytes, little-endian
                                       little-endian
ad17835b00000000 - Timestamp, 8
00000000000000000 - Network servi
                                    of receiver, 8 bytes, little-endian
000000000000000000000ffff000000
                                  Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of recei, 2 bytes, 8333
000000000000000 - Network rvices of sender, 8 bytes, little-endian
00000000000000000000ffff0 0000 - Network address of sender, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of Jender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

000000000000ffff0

6a8d7a440ec27a11b

What software am I using?

– var str

Version

```
2f70726f6772616d6d696e67626c6f6
                                                       J00001
7f110100 - Protocol version, 4 bytes, little-ex
                                                  70015
0000000000000000 - Network services of sender, ₹
                                                     little-endian
ad17835b00000000 - Timestamp, 8 bytes, little-end
0000000000000000 - Network services of receiver, 8
                                                     , little-endian
ceiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of receiver, 2 bytes, 8333
0000000000000000 - Network services of sender, 8 bytes,
                                                       ttle-endian
000000000000000000000ffff00000000 - Network address of sen
                                                        r. 16 bytes. IPv4
        0.0.0.0
8d20 - Network port of sender, 2 bytes, 8333
f6a8d7a440ec27a1 - Nonce, 8 bytes, used for communicating responses
1b2f70726f6772616d6d696e67626c6f636b636861696e3a302e312f - User agent
        /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

00000008d200000000000000000

000000000000ffff0

6a8d7a440ec27a11b

00001

The last block that the node sending the message has!

Version

```
7f110100 - Protocol version, 4 by
                                  tle-endian, 70015
00000000000000000 - Network servi
                                 sender, 8 bytes, little-endian
ad17835b000000000 - Timestamp, 8
                                , little-endian
00000000000000000 - Network se
                               of receiver, 8 bytes, little-endian
Network address of receiver, 16 bytes, IPv4
       0.0.0.0
8d20 - Network port of refer, 2 bytes, 8333
000000000000000 - Networkservices of sender, 8 bytes, little-endian
0.0.0.0
8d20 - Network port f sender, 2 bytes, 8333
f6a8d7a440ec27a1 Monce, 8 bytes, used for communicating responses
1b2f70726f677261 16d696e67626c6f636b636861696e3a302e312f - User agent
       /programmingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

00000008d200000000000000000

2f70726f6772616d6d696e67626c6f6

0000000000000ffff0 6a8d7a440ec27a11b

Bloom filters!

Version

```
2f70726f6772616d6d696e67626c6f6
                                                         00001
7f110100 - Protocol version, 4 by
                                       Ztle-endian, 70015
00000000000000000 - Network servi
                                      sender, 8 bytes, little-endian
                                   5. little-endian
ad17835b000000000 - Timestamp,
00000000000000000 - Network se
                                   of receiver, 8 bytes, little-endian
0000000000000000000000000ffff000
                                  Network address of receiver, 16 bytes, IPv4
        0.0.0.0
8d20 - Network port of r
                          er, 2 bytes, 8333
0000000000000000 - Netw
                          services of sender, 8 bytes, little-endian
                        70000000 - Network address of sender, 16 bytes, IPv4
0.0.0.0
8d20 - Network por of sender, 2 bytes, 8333
f6a8d7a440ec27a1 Nonce, 8 bytes, used for communicating responses
1b2f70726f67726 3d6d696e67626c6f636b636861696e3a302e312f - User agent
        /prog/ammingblockchain:0.1/
00000000 - Height, 0
01 - Optional flag for relay, based on BIP37
```

00000008d2000000000000000000

Version

```
class VersionMessage:
    command = b'version'
   def init (self, version=70015, services=0, timestamp=None,
                receiver_services=0,
                receiver ip=b'\x00\x00\x00', receiver port=8333,
                sender services=0,
                 sender ip=b'\x00\x00\x00\x00', sender_port=8333,
                 nonce=None, user_agent=b'/programmingbitcoin:0.1/',
                 latest block=0, relay=False):
       self.version = version
       self.services = services
       if timestamp is None:
           self.timestamp = int(time.time())
           self.timestamp = timestamp
       self.receiver services = receiver services
       self.receiver ip = receiver ip
       self.receiver port = receiver port
       self.sender_services = sender_services
       self.sender ip = sender ip
       self.sender port = sender port
       if nonce is None:
           self.nonce = int to little endian(randint(0, 2**64), 8)
           self.nonce = nonce
        self.user agent = user agent
       self.latest block = latest block
        self.relav = relav
```

```
class VersionMessage:
    command = b'ver[
                       def serialize(self):
    def init (se
                           result = int_to_little_endian(self.version, 4)
                           result += int_to_little_endian(self.services, 8)
                   se
se
no
                           result += int_to_little_endian(self.timestamp, 8)
                           result += int to little endian(self.receiver services, 8)
         self.versio
                           # IPV4 is 10 00 bytes and 2 ff bytes then receiver ip
                           result += b'\x00' * 10 + b'\xff\xff' + self.receiver_ip
         self.servic
         if timestam
                           result += self.receiver port.to bytes(2, 'big')
             self.ti
                           result += int to little endian(self.sender services, 8)
             self.ti
                           # IPV4 is 10 00 bytes and 2 ff bytes then sender ip
                           result += b'\x00' * 10 + b'\xff\xff' + self.sender_ip
         self.receiv
         self.receiv
                           result += self.sender port.to bytes(2, 'big')
         self.receiv
                           # nonce should be 8 bytes
         self.sender
                           result += self.nonce
         self.sender
         self.sender
                           result += encode_varint(len(self.user_agent))
         if nonce is
                           result += self.user_agent
             self.no
                           result += int_to_little_endian(self.latest_block, 4)
             self.no
                           if self.relav:
         self.user_a
                               result += b'\x01'
         self.latest
                               result += b'\x00'
         self.relay
                           return result
```

Version

We will only send!

Verack

verack

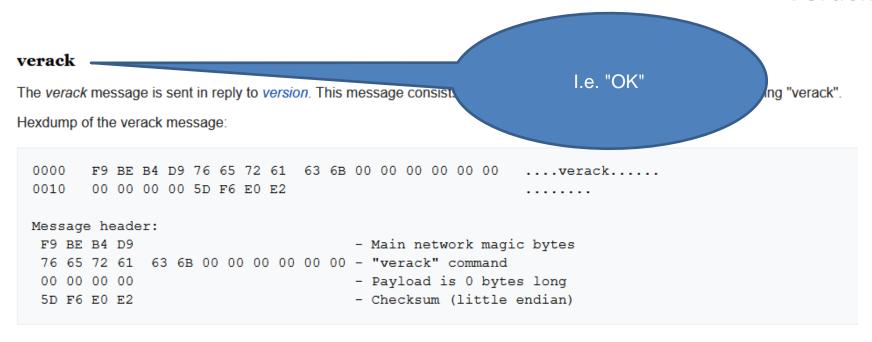
The *verack* message is sent in reply to *version*. This message consists of only a message header with the command string "verack". Hexdump of the verack message:

```
0000 F9 BE B4 D9 76 65 72 61 63 6B 00 00 00 00 00 ....verack.....

0010 00 00 00 5D F6 E0 E2 .......

Message header:
F9 BE B4 D9 - Main network magic bytes
76 65 72 61 63 6B 00 00 00 00 0 - "verack" command
00 00 00 00 - Payload is 0 bytes long
5D F6 E0 E2 - Checksum (little endian)
```

Verack



Verack

```
class VerAckMessage:
    command = b'verack'
    def __init__(self):
        pass
    @classmethod
    def parse(cls, s):
        return cls()
    def serialize(self):
        return b''
```

```
class SimpleNode:
    def __init__(self, host, port=None, testnet=False, logging=False):
        if port is None:
            if testnet:
                port = 18333
            else:
                port = 8333
        self.testnet = testnet
        self.logging = logging
        # connect to socket
        self.socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        self.socket.connect((host, port))
        # create a stream that we can use with the rest of the library
        self.stream = self.socket.makefile('rb', None)
```

```
class SimpleNode:
def send(self, message):
    '''Send a message to the connected node'''
    # create a network envelope
    envelope = NetworkEnvelope(
        message.command, message.serialize(), testnet=self.testnet)
    if self.logging:
        print('sending: {}'.format(envelope))
    # send the serialized envelope over the socket using sendall
    self.socket.sendall(envelope.serialize())
```

```
class SimpleNode:
      def __init__(self, host, port=None, testnet=False, logging=False):
def read(self):
    '''Read a message from the socket'''
    envelope = NetworkEnvelope.parse(self.stream, testnet=self.testnet)
    if self.logging:
         print('receiving: {}'.format(envelope))
    return envelope
         self.socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
         self.socket.connect((host, port))
         # create a stream that we can use with the rest of the library
         self.stream = self.socket.makefile('rb', None)
```

```
class SimpleNode:
      def wait_for(self, *message_classes):
                                                                                 alse):
          command = None
          command_to_class = {m.command: m for m in message_classes}
          while command not in command_to_class.keys():
              # get the next network message
              envelope = self.read()
              command = envelope.command
                                                                                 STREAM)
              if command == VersionMessage.command:
                  # send verack
                  self.send(VerAckMessage())
                                                                                 ibrary
              elif command == PingMessage.command:
                  self.send(PongMessage(envelope.payload))
          # return the envelope parsed as a member of the right message class
          return command to class[command].parse(envelope.stream())
```

class PingMessage: command = b'ping' def __init__(self, nonce): self.nonce = nonce @classmethod def parse(cls, s): nonce = s.read(8) return cls(nonce) def serialize(self): return self.nonce class PongMessage: command = b'pong' def __init__(self, nonce): self.nonce = nonce def parse(cls, s): nonce = s.read(8) return cls(nonce) def serialize(self): return self.nonce

A simple node

Ping pong

```
To simplify things, our
                                                                          implementation will be
class SimpleNode:
                                                                               sinchronous!
      def wait_for(self, *message_classes):
                                                                      (i.e. I send a message and wait
                                                                                for a reply)
          # initialize the command we have, which should be None
          command = None
          command_to_class = {m.command: m for m in message_classes}
          while command not in command to class.keys():
              envelope = self.read()
              command = envelope.command
                                                                                STREAM)
              if command == VersionMessage.command:
                  # send verack
                  self.send(VerAckMessage())
                                                                                ibrarv
              elif command == PingMessage.command:
                  self.send(PongMessage(envelope.payload))
          # return the envelope parsed as a member of the right message class
          return command to class[command].parse(envelope.stream())
```

```
class SimpleNode:
      def wait_for(self, *message_classes):
                                                                                 alse):
          command = None
          command_to_class = {m.command: m for m in message_classes}
          while command not in command_to_class.keys():
              # get the next network message
              envelope = self.read()
              command = envelope.command
                                                                                 STREAM)
              if command == VersionMessage.command:
                  # send verack
                  self.send(VerAckMessage())
                                                                                 ibrary
              elif command == PingMessage.command:
                  self.send(PongMessage(envelope.payload))
          # return the envelope parsed as a member of the right message class
          return command to class[command].parse(envelope.stream())
```

```
To simplify things, our
                                                                          implementation will be
class SimpleNode:
                                                                               sinchronous!
      def wait_for(self, *message_classes):
                                                                      (i.e. I send a message and wait
                                                                                for a reply)
          # initialize the command we have, which should be None
          command = None
          command_to_class = {m.command: m for m in message_classes}
          while command not in command to class.keys():
              envelope = self.read()
              command = envelope.command
                                                                                STREAM)
              if command == VersionMessage.command:
                  # send verack
                  self.send(VerAckMessage())
                                                                                ibrarv
              elif command == PingMessage.command:
                  self.send(PongMessage(envelope.payload))
          # return the envelope parsed as a member of the right message class
          return command to class[command].parse(envelope.stream())
```

To simplify things, our

```
implementation will be
class SimpleNode:
                                                                               sinchronous!
      def wait_for(self, *message_classes):
                                                                     (i.e. I send a message and wait
                                                                                for a reply)
          # initialize the command we have, which should be None
          command = None
          command_to_class = {m.command: m for m in message_classes
          while command not in command_to_class.keys():
              envelope = self.read()
              command = envelope.command
                                                                 Real nodes communicate in
              if command == VersionMessage.command:
                                                                   asynchronous manner!
                  # send verack
                  self.send(VerAckMessage())
              elif command == PingMessage.command:
                  self.send(PongMessage(envelope.payload))
          # return the envelope parsed as a member of the right message class
          return command to class[command].parse(envelope.stream())
```

```
class SimpleNode:
def handshake(self):
    '''Do a handshake with the other node.
    Handshake is sending a version message and getting a verack back.'''
    # create a version message
    version = VersionMessage()
    # send the command
    self.send(version)
    # wait for a verack message
    self.wait_for(VerAckMessage)
```

```
class SimpleNode:
def handshake(self):
    '''Do a handshake with the other node.
    Handshake is sending a version
                                        rage and getting a verack back.'''
    # create a version message
    version = VersionMessage()
    # send the command
    self.send(version)
                                                 This is how we start
    # wait for a verack message
    self.wait for(VerAckMessage)
```

Initial connection

```
# Establish a connection to a testnet node
node = SimpleNode('testnet.programmingbitcoin.com', testnet=True)
node.handshake()
```

I need to know the address of at lesat one node!

SPV nodes

If I am an SPV node, I need the blockchain of headers!

- Version
- Verack
- GetHeaders
- Headers

GetHeaders

With getheaders I can obtain headers (max 2000!)

Payload:

- 7f110100 Protocol version, 4 bytes, little-endian, 70015
- 01 Number of hashes, varint
- a35b...00 Starting block, little-endian
- 0000...00 Ending block, little-endian

GetHeaders

0000000000000

With **getheaders** I can obtain headers (max

7f110100**01**a35bd0ca2f4a88c4eda6d213e2378a575

Payload:

0000000000

- 7f110100 Protocol version, 4 bytes, little-endian, 70015
- 01 Number of hashes, varint
- a35b...00 Starting block, little-endian
- 0000...00 Ending block, little-endian

In reality, I can start from multiple blocks (e.g. for forks)

GetHeaders

With getheaders I can obtain headers (ma

If this is 0, we should receive **up to** the next 2000 blocks

Payload:

7f110100**01**a35bd0ca2f4a88c4eda6d213e2378a5758dfcd6a7

- 7f110100 Protocol version, 4 by , lity
- 01 Number of hashes, varint
- a35b...00 Starting block, little-endian
- 0000...00 Ending block, little-endian

00000000

Depends on the node we connect to!

```
class GetHeadersMessage:
    command = b'getheaders'
    def __init__(self, version=70015, num hashes=1, start block=None, end block=None):
        self.version = version
        self.num_hashes = num_hashes
        if start block is None:
            raise RuntimeError('a start block is required')
        self.start block = start block
        if end block is None:
            self.end_block = b' \times 200' * 32
            self.end block = end block
    def serialize(self):
        '''Serialize this message to send over the network'''
        # protocol version is 4 bytes little-endian
        result = int to little endian(self.version, 4)
        # number of hashes is a varint
        result += encode_varint(self.num_hashes)
        result += self.start_block[::-1]
        # end block is also in little-endian
        result += self.end_block[::-1]
        return result
```

GetHeaders

headers

Reply to getheaders is headers

Payload:

0200000020df3b053dc46f162a9b00c7f0d5124e2676d47bbe7c5d0793a50000000000000000ef445fef 2ed495c275892206ca533e7411907971013ab83e3b47bd0d692d14d4dc7c835b67d8001ac157e67**00** 00000002030eb2540c41025690160a1014c577061596e32e426b712c7ca00000000000000768b89f07 044e6130ead292a3f51951adbd2202df447d98789339937fd006bd44880835b67d8001ade092046**00**

- 02 Number of block headers
- 00...67 Block header
- 00 Number of transactions (always 0)

headers

Reply to getheaders is headers

We have to *parse* this!!!

Payload:

0200000020df3b053dc46f162a9b00c7f0d5124e2 47bbe7c5d0793a50000000000000000ef445fef 2ed495c275892206ca533e7411907971013ab8 47bd0d692d14d4dc7c835b67d8001ac157e67**00** 00000002030eb2540c41025690160a1014c 061596e32e426b712c7ca000000000000000768b89f07 044e6130ead292a3f51951adbd2202df d98789339937fd006bd44880835b67d8001ade092046**00**

- 02 Number of block headers
- 00...67 Block header
- 00 Number of transactions (always 0)

Block Headers

Block headers are sent in a headers packet in response to a getheaders message.

Field Size	Description	Data type	Comments
4	version	int32_t	Block version information (note, this is signed)
32	prev_block	char[32]	The hash value of the previous block this particular block references
32	merkle_root	char[32]	The reference to a Merkle tree collection which is a hash of all transactions related to this block
4	timestamp	uint32_t	A timestamp recording when this block was created (Will overflow in 2106 ^[2])
4	bits	uint32_t	The calculated difficulty target being used for this block
4	nonce	uint32_t	The nonce used to generate this block to allow variations of the header and compute different hashes
1+	txn_count	var_int	Number of transaction entries, this value is always 0

https://en.bitcoin.it/wiki/Protocol_documentation

Our implementation checks this in the payload of the header message!

Block Headers

Block headers are sent in a headers packet in response to a getheaders message.

Field Size	Description	Data type	omments
4	version	int32_t	Block version information (possessigned)
32	prev_block	char[32]	The hash value of the lock this particular block references
32	merkle_root	char[32]	The reference a Merkle tree collection which is a hash of all transactions related to this block
4	timestamp	uint32_t	A time and the state of the sta
4	bits	uint32	The calculated difficulty target being used for this block
4	nonce	uint32_t	The nonce used to generate this block to allow variations of the header and compute different hashes
1+	txn_count	var_int	Number of transaction entries, this value is always 0

In bytes

020000208ec39428b17323fa0ddec8e887b4a7c53b8c0a0 a220cfd00000000000000005b0750fce0a889502d4050 8d39576821155e9c9e3f5c3157f961db38fd8b25be1e77a 759e93c0118a4ffd71d

- 02000020 version, 4 bytes, LE
- 8ec3...00 previous block, 32 bytes, LE
- 5b07...be merkle root, 32 bytes, LE
- 1e77a759 timestamp, 4 bytes, LE
- e93c0118 bits, 4 bytes
- a4ffd71d nonce, 4 bytes

In bytes

020000208ec39428b17323fa0ddec8e887b4a7c53b8c0a0 a220cfd00000000000000005b0750fce0a889502d4050 8d39576821155e9c9e3f5c3157f961db38fd8b25be1e77a 759e93c0118a4ffd71d

- 02000020 version, 4 bytes, LE
- 8ec3...00 previous block, 32 bytes, LE
- 5b07...be merkle root, 32 bytes, LE
- 1e77a759 timestamp, 4 bytes, LE
- e93c0118 bits, 4 bytes
- a4ffd71d nonce, 4 bytes

In Python

```
class Block:
    def __init__(self, version, prev_block, merkle_root,
                 timestamp, bits, nonce):
        self.version = version
        self.prev_block = prev_block
        self.merkle_root = merkle_root
        self.timestamp = timestamp
        self.bits = bits
        self.nonce = nonce
```

In Python

```
class Block:
     @classmethod
     def parse(cls, s):
         '''Takes a byte stream and parses a block. Returns a Block object'''
         # s.read(n) will read n bytes from the stream
         # version - 4 bytes, little endian, interpret as int
         version = little endian to int(s.read(4))
         prev_block = s.read(32)[::-1]
         # merkle_root - 32 bytes, little endian (use [::-1] to reverse)
         merkle_root = s.read(32)[::-1]
         # timestamp - 4 bytes, little endian, interpret as int
         timestamp = little_endian_to_int(s.read(4))
         # bits - 4 bytes
         bits = s.read(4)
         # nonce - 4 bytes
         nonce = s.read(4)
         # initialize class
         return cls(version, prev_block, merkle_root, timestamp, bits, nonce)
```

https://github.com/jimmysong/programmingbitcoin/blob/master/ch09.asciidoc

In Python

```
class Block:
     def serialize(self):
                                                              ot,
         '''Returns the 80 byte block header'''
         # version - 4 bytes, little endian
         result = int to little endian(self.version, 4)
         # prev block - 32 bytes, little endian
         result += self.prev_block[::-1]
         # merkle_root - 32 bytes, little endian
         result += self.merkle_root[::-1]
         # timestamp - 4 bytes, little endian
         result += int to little endian(self.timestamp, 4)
         # bits - 4 bytes
         result += self.bits
         # nonce - 4 bytes
         result += self.nonce
         return result
```

https://github.com/jimmysong/programmingbitcoin/blob/master/ch09.asciidoc

Now we know how to

headers

Reply to getheaders is headers

parse this part!!!

Payload:

0200000020df3b053dc46f162a9b00c7f0d5124e2 47bbe7c5d0793a50000000000000000ef445fef 2ed495c275892206ca533e7411907971013ab8 47bd0d692d14d4dc7c835b67d8001ac157e67**00** 00000002030eb2540c41025690160a1014c 061596e32e426b712c7ca000000000000000768b89f07 044e6130ead292a3f51951adbd2202df d98789339937fd006bd44880835b67d8001ade092046**00**

- 02 Number of block headers
- 00...67 Block header
- 00 Number of transactions (always 0)

```
class HeadersMessage:
    command = b'headers'
    def __init__(self, blocks):
        self.blocks = blocks
    @classmethod
    def parse(cls, stream):
        # number of headers is in a varint
        num headers = read varint(stream)
        # initialize the blocks array
        blocks = []
        # loop through number of headers times
        for _ in range(num_headers):
            # add a block to the blocks array by parsing the stream
            blocks.append(Block.parse(stream))
            # read the next varint (num txs)
            num_txs = read_varint(stream)
            # num txs should be 0 or raise a RuntimeError
            if num txs != 0:
                raise RuntimeError('number of txs not 0')
        # return a class instance
        return cls(blocks)
```

```
class HeadersMessage:
    command = b'headers'
    def __init__(self, blocks):
        self.blocks = blocks
                                                                           Blockchain of headers
    @classmethod
    def parse(cls, stream):
        # number of headers is in a varint
        num headers = read varint(stream)
        # initialize the blocks array
        blocks = []
        # loop through number of headers times
        for _ in range(num_headers):
            # add a block to the blocks array by parsing the stream
            blocks.append(Block.parse(stream))
            # read the next varint (num txs)
            num_txs = read_varint(stream)
            if num txs != 0:
                raise RuntimeError('number of txs not 0')
        # return a class instance
        return cls(blocks)
```

```
class HeadersMessage:
    command = b'headers'
    def __init__(self, blocks):
        self.blocks = blocks
                                                                            Number of headers
    @classmethod
    def parse(cls, stream):
        # number of headers is in a varint
        num headers = read varint(stream)
        # initialize the blocks array
        blocks = []
        # loop through number of headers times
        for _ in range(num_headers):
            # add a block to the blocks array by parsing the stream
            blocks.append(Block.parse(stream))
            # read the next varint (num txs)
            num_txs = read_varint(stream)
            if num txs != 0:
                raise RuntimeError('number of txs not 0')
        # return a class instance
        return cls(blocks)
```

We will get headers message

```
class HeadersMessage:
    command = b'headers'
    def __init__(self, blocks):
        self.blocks = blocks
    @classmethod
    def parse(cls, stream):
        num headers = read varint(stream)
        blocks = []
        # loop through number of headers times
        for _ in range(num_headers):
            # add a block to the blocks array by plang the stream
            blocks.append(Block.parse(stream))
            # read the next varint (num txs)
            num_txs = read_varint(stream)
            if num txs != 0:
                raise RuntimeError('number of txs not 0')
        # return a class instance
        return cls(blocks)
```

Headers (we will store this as objects of the class block)

```
class HeadersMessage:
    command = b'headers'
    def __init__(self, blocks):
        self.blocks = blocks
                                                                    Number of transactions needs to be
    @classmethod
                                                                              0 (for a header)!
    def parse(cls, stream):
        num headers = read varint(stream)
        blocks = []
        # loop through number of headers times
        for _ in range(num_headers):
            # add a block to the blocks array by parsing
            blocks.append(Block.parse(stream))
            # read the next varint (num txs)
            num_txs = read_varint(stream)
            # num txs should be 0 or raise a RuralmeError
            if num txs != 0:
                raise RuntimeError('number of txs not 0')
        # return a class instance
        return cls(blocks)
```

```
# Establish a connection to a testnet node
node = SimpleNode('testnet.programmingbitcoin.com', testnet=True)
node.handshake()
# Get the first 2000 blocks of Bitcoin:
previous = Block.parse(BytesIO(GENESIS_BLOCK))
getheaders = GetHeadersMessage(start_block = previous.hash())
node.send(getheaders)
headers = node.wait for(HeadersMessage)
for x in headers.blocks:
    print(x.hash().hex(),x.prev_block.hex())
```

```
# Establish a connection to a testnet
node = SimpleNode('testnet.programming)
                                                We need the genesis block!
node.handshake()
# Get the first 2000 blocks of Bitcoin:
previous = Block.parse(BytesIO(GENESIS_BLOCK))
getheaders = GetHeadersMessage(start_block = previous.hash())
node.send(getheaders)
headers = node.wait_for(HeadersMessage)
for x in headers.blocks:
    print(x.hash().hex(),x.prev_block.hex())
```

```
# Establish a connection to a testnet
node = SimpleNode('testnet.programming)
                                                    Ask for blocks!
node.handshake()
# Get the first 2000 blocks of
previous = Block.parse(Byte GENESIS_BLOCK))
getheaders = GetHeader_message(start_block = previous.hash())
node.send(getheaders)
headers = node.wait_for(HeadersMessage)
for x in headers.blocks:
    print(x.hash().hex(),x.prev_block.hex())
```

```
# Establish a connection to a testnet
node = SimpleNode('testnet.programming)
                                           Wait to recieve the blocks
node.handshake()
# Get the first 2000 blocks of Bitco
getheaders = GetHeadersMessa (start_block = previous.hash())
node.send(getheaders)
headers = node.wait for(HeadersMessage)
for x in headers.blocks:
   print(x.hash().hex(),x.prev_block.hex())
```

Allos to check if POW is valid!

Block Headers

Block headers are sent in a headers

(do not confuse with bits of a Merkle tree)

Field Size	Description	Data type	enenes
4	version	int32_t	Block version (note, this is signed)
32	prev_block	char[32]	The hash the previous block this particular block references
32	merkle_root	char[32]	The ref ence to a Merkle tree collection which is a hash of all transactions related to this block
4	timestamp	uint32_t	A timestamp recording when this block was created (Will overflow in 2106 ^[2])
4	bits	uint32_t	The calculated difficulty target being used for this block
4	nonce	uint32_t	The nonce used to generate this block to allow variations of the header and compute different hashes
1+	txn_count	var_int	Number of transaction entries, this value is always 0

How to check that POW of a block is valid?

- hash(header) < target
- 2. bits are correct (according to the difficulty adjustment)

How to check that POW of a block is valid?

- hash(header) < target
- 2. bits are correct (according the difficulty adjustment)

To compute the hash of a block I only need the header (it conatins the MerkleRoot)

How to check that POW of a block is valid?

- 1. hash(header) < target
- bits are correct (according to the difficulty adjustment)

Difficulty adjustment is deterministic!

All the nodes can check it!

Here we use the timestamp of a block!

How to check that POW of a block is valid?

- hash(header) < target
- 2. bits are correct (according the difficulty adjustment)

Target is computed using bits (the ones from the block header)

Target

How to computetarget?

Usando bits = 4 bytes

bits are two different numbers:

- 1. Byte 4 is the exponent: exp
- 2. Bytes 1,2,3 are the coefficient: coeff (in Little-endian)

target = $coeff \cdot 256^{exp-3}$

Target

```
target = coeff \cdot 256^{exp-3}
```

```
def bits_to_target(bits):
    '''Turns bits into a target (large 256-bit integer)'''
    # last byte is exponent
    exponent = bits[-1]
    # the first three bytes are the coefficient in little endian
    coefficient = little_endian_to_int(bits[:-1])
    # the formula is:
    # coefficient * 256**(exponent-3)
    return coefficient * 256**(exponent - 3)
```

How to check that POW of a block is valid?

- 1. hash(header) < target
- 2. bits are correct (according to the difficulty adjustment)

Block header already has everything needed!!!

```
def serialize(self):
    '''Returns the 80 byte block header'''
    # version - 4 bytes, little endian
    result = int to little endian(self.version, 4)
    # prev block - 32 bytes, little endian
    result += self.prev block[::-1]
    # merkle root - 32 bytes, little endian
    result += self.merkle_root[::-1]
    # timestamp - 4 bytes, little endian
    result += int_to_little_endian(self.timestamp, 4)
    # bits - 4 bytes
    result += self.bits
    # nonce - 4 bytes
    result += self.nonce
    return result
```

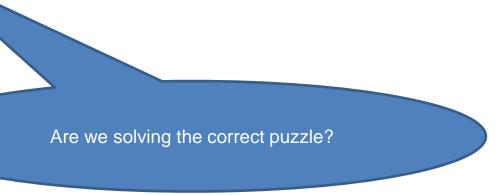
class Block:

```
class Block:
    def hash(self):
         '''Returns the hash256 interpreted little endian of the block'''
         # serialize
         s = self.serialize()
         # hash256
        h256 = hash256(s)
         # reverse
         return h256[::-1]
```

```
class Block:
     def check_pow(self):
         '''Returns whether this block satisfies proof of work'''
         # get the hash256 of the serialization of this block
         h256 = hash256(self.serialize())
         # interpret this hash as a little-endian number
         proof = little_endian_to_int(h256)
         # return whether this integer is less than the target
         return proof < self.target()</pre>
```

How to check that POW of a block is valid?

- 1. hash(header) < target
- 2. bits are correct (according to the difficulty adjustment)



Difficulty adjustment

bits are correct (according the difficulty adjustment)

- Each group of 2016 blocks in Bitcoin is called difficulty adjustment period
- Target is adjusted according to the formula:

```
time_diff = (timestamp of the last block) - (timestamp of the first block)
```

```
new_target = prev_target ·time_diff/(2 weeks)
```

ficulty adjustment

E.g. for the period of blocks: 2017 until 4032 new_target is computed using the difference between blocks 2016 and 1

bits are correct

- Each group of 2019
- Target is adjusted

anniculty adjustment period

ang to the formula:

time_diff = (timestamp of the last block) - (timestamp of the first block)

new_target = prev_target ·time_diff/(2 weeks)

annoulty adjustment period

ficulty adjustment

E.g. for the period of blocks: 2017 until 4032 This is the target fo the period 1 - 2016

bits are correct

- Each group of 2010 >
- Target is adjusted a to the formula:

```
time_diff = (timestar of the last block) – (timestamp of the first block)
```

```
new_target = prev_target ·time_diff/(2 weeks)
```

annoulty adjustment period

sticulty adjustment

Remark:

If time_diff>8 weeks we will use 8 weeks If time_diff<3.5 days we will use 3.5 days

bits are correct

- Each group of 201y
- Target is adjusted along to the formula:

time_diff = (timestamp of the last block) - (timestamp of the first block)

new_target = prev_target ·time_diff/(2 weeks)

Siculty adjustment

Remark:

Another one-off bug from Satoshi: the difference is between 2015 blocks, so not precisely 2 weeks!

bits are corre

- Each group of 201y
- Target is adjusted

emiculty adjustment period

aing to the formula:

time_diff = (timestamp of the last block) - (timestamp of the first block)

new_target = prev_target ·time_diff/(2 weeks)

Difficulty adjustment

```
def calculate_new_bits(previous_bits, time_differential):
    '''Calculates the new bits given
    a 2016-block time differential and the previous bits'''
    # if the time differential is greater than 8 weeks, set to 8 weeks
    if time differential > TWO WEEKS * 4:
        time differential = TWO WEEKS * 4
    # if the time differential is less than half a week, set to half a week
    if time_differential < TWO_WEEKS // 4:</pre>
        time_differential = TWO_WEEKS // 4
    # the new target is the previous target * time differential / two weeks
    new_target = bits_to_target(previous_bits) * time_differential // TWO_WEEKS
    # if the new target is bigger than MAX_TARGET, set to MAX_TARGET
    if new_target > MAX_TARGET:
        new target = MAX TARGET
    # convert the new target to bits
    return target to bits(new target)
```

Function implemented in helper.py

ficulty adjustment Target of the genesys block; it can never be lower than this! def cald a 2016-bloc. # if the time differen weeks, set to 8 weeks if time_differential time_differential ∠KS * # if the time differe less than half a week, set to half a week if time_differential WEEKS // 4: time differentia WO WEEKS // 4 # the new target is previous target * time differential / two weeks new_target = bits_t// carget(previous_bits) * time_differential // TWO_WEEKS # if the new target is bigger than MAX_TARGET, set to MAX_TARGET if new_target > MAX_TARGET:

Function implemented in helper.py

new_target = MAX_TARGET
convert the new target to bits
return target_to_bits(new_target)

Difficulty adjustment

```
def target_to_bits(target):
    '''Turns a target integer back into bits, which is 4 bytes'''
    raw_bytes = target.to_bytes(32, 'big')
    raw_bytes = raw_bytes.lstrip(b'\x00')
   if raw_bytes[0] > 0x7f:
        # if the first bit is 1, we have to start with 00
        exponent = len(raw_bytes) + 1
        coefficient = b'\x00' + raw_bytes[:2]
   else:
        # otherwise, we can show the first 3 bytes
        # exponent is the number of digits in base-256
        exponent = len(raw bytes)
       # coefficient is the first 3 digits of the base-256 number
        coefficient = raw_bytes[:3]
   # we've truncated the number after the first 3 digits of base-256
    new_bits = coefficient[::-1] + bytes([exponent])
   return new_bits
```

With this, our implementation can download the block headers, and verify that they all satisfy the proof-of-work for this block!

As an exercise, you can use our implementation to verify this!

References

You can read more in:

- Programming Bitcoin, chapters 9,10
- https://en.bitcoin.it/wiki/Protocol_documentation