

# The Bitcoin Network

How Bitcoin works?

# Bitcoin network

## **A standard p2p network:**

- Nodes connect and disconnect
- New nodes enter the network
- There is no perfect connectivity
- Nodes listen to the messages by their neighbours
- Nodes transmit messages to their neighbours

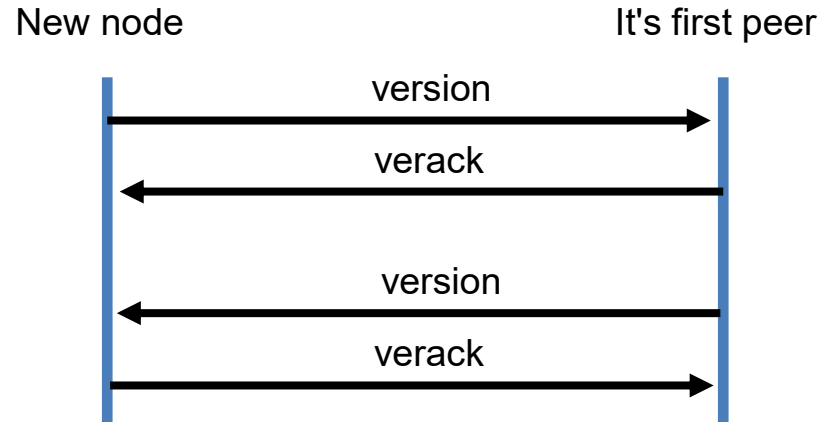
## **How does a new node get to have neighbours?**

- Static DNS Seed
- DNS Seed BIND (Berkeley Internet Name Daemon)
- A specific IP of a node already known to the new node

# Bitcoin network

## Basic messages to establish a connection:

- *version*
- *verack*

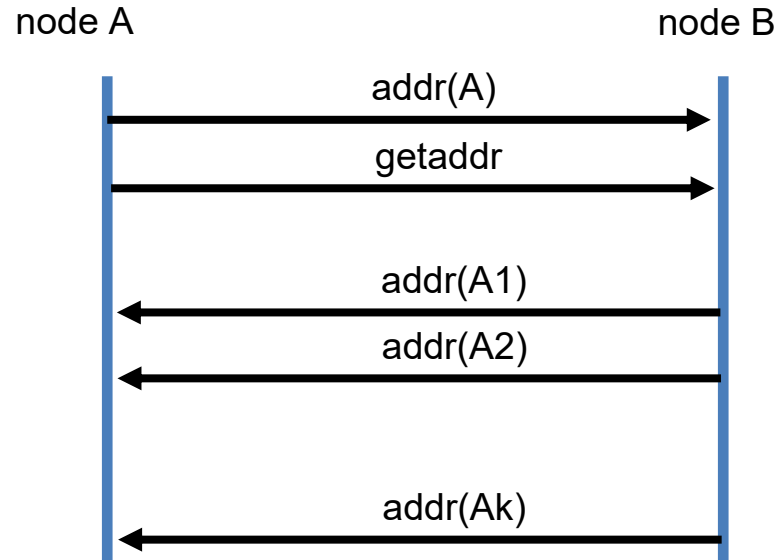


## **Discovering new peers:**

- A connects to B
- B will pass on the address of A to its peers
- A can also ask B for the addresses of its peers
- B replies with these addresses

# Bitcoin network

## Discovering new peers:



## **Nodes enter and leave the network:**

- A dynamic process
- Node stores its peers and upon entering again tries to connect to them
- If there are no messages from a node for 30 min, a message is sent
- If there are no messages for 90 min, the node is disconnected (one less edge)
  
- It is recommended to have no more than 100 peers
- Too many connections put pressure onto the network

# Bitcoin network

## Basic communication:


tx → node A

node B




# Bitcoin network

## Basic communication:




# Bitcoin network

## Basic communication:



# Bitcoin network

## Basic communication:



# Bitcoin network


**Basic communication (blocks):**



node B


# Bitcoin network

Basic communication (blocks):



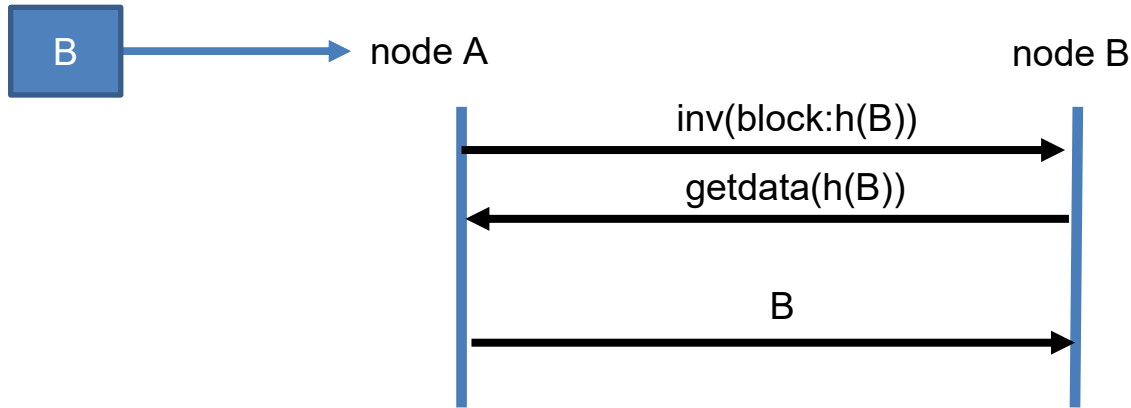
# Bitcoin network

## Basic communication (blocks):



# Bitcoin network

## Basic communication (blocks):



# Bitcoin network

## **Full nodes:**

- Store the entire blockchain
- Verify all the transactions and blocks
- How does a full node obtain the blockchain upon entering the network?



# Bitcoin network

Up to version 0.9.3 of BitcoinCore

## How does node X obtain the blockchain?

- Always starts with the genesis block (hardcoded in its software)
- Other blocks are received by its peers on the network
- First message sent/received by X: *version* (contains the field *BestHeight*)
- X sends *getblocks* to its peer with the highest *BestHeight* (call this node Y)
- The message *getblocks* has the hash of the last block in the blockchain of X
- Y identifies the first 500 blocks that X does not yet have
- Y sends the message *inv* with the hashes of these 500 blocks
- To obtain a block, X sends to Y *getdata(hash)* message
- Maximum of 500 *getdata* can be active per connection

# Initial block download

node A

node B

# Initial block download

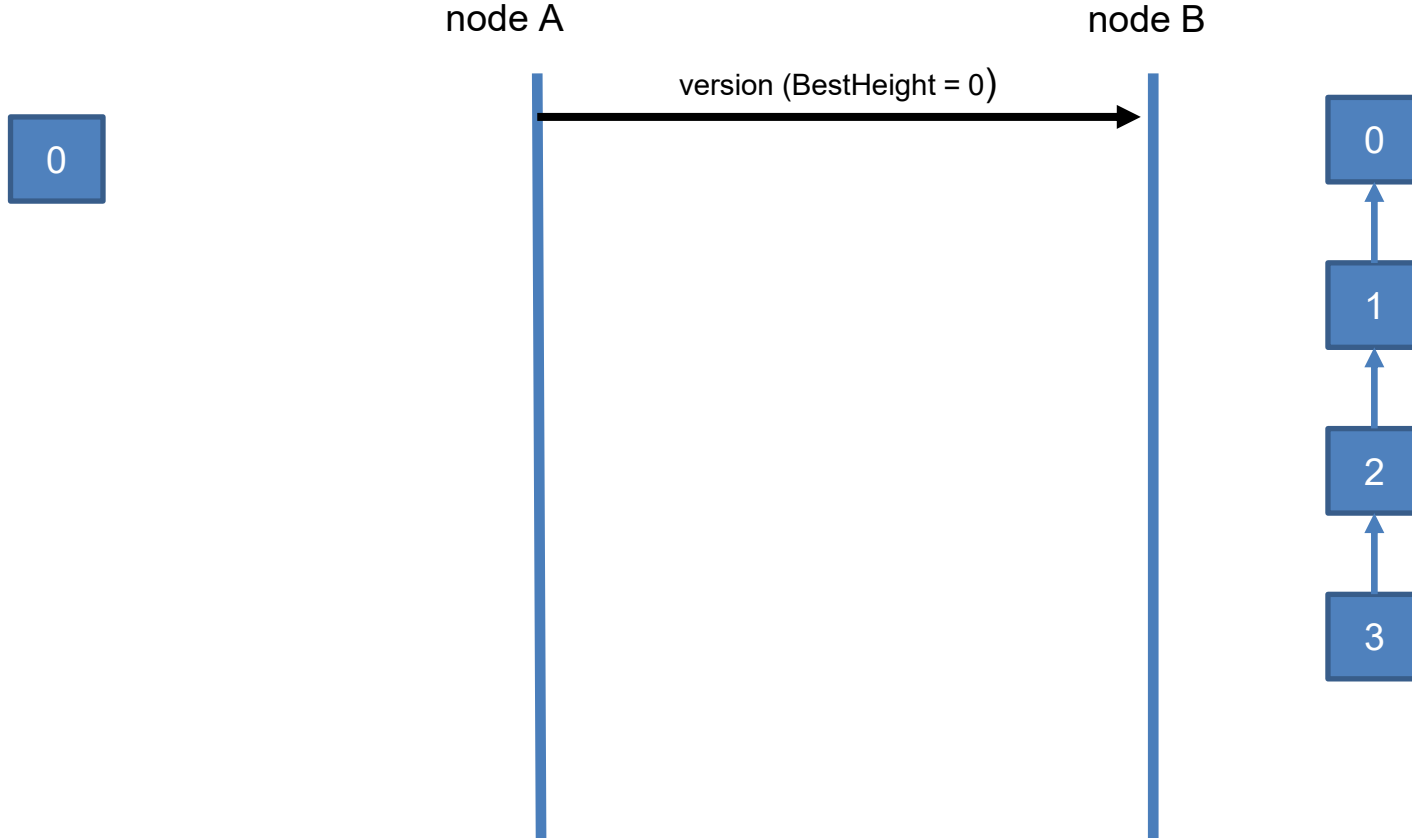
node A



node B



# Initial block download



# Initial block download

0

node A

node B

version (BestHeight = 0)

version (BestHeight = 3)

0

1

2

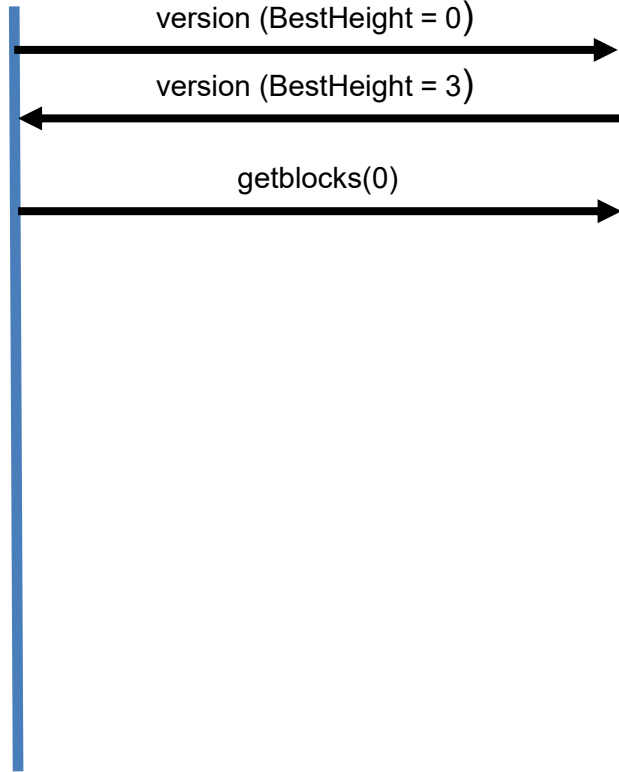
3

# Initial block download

0

node A

node B

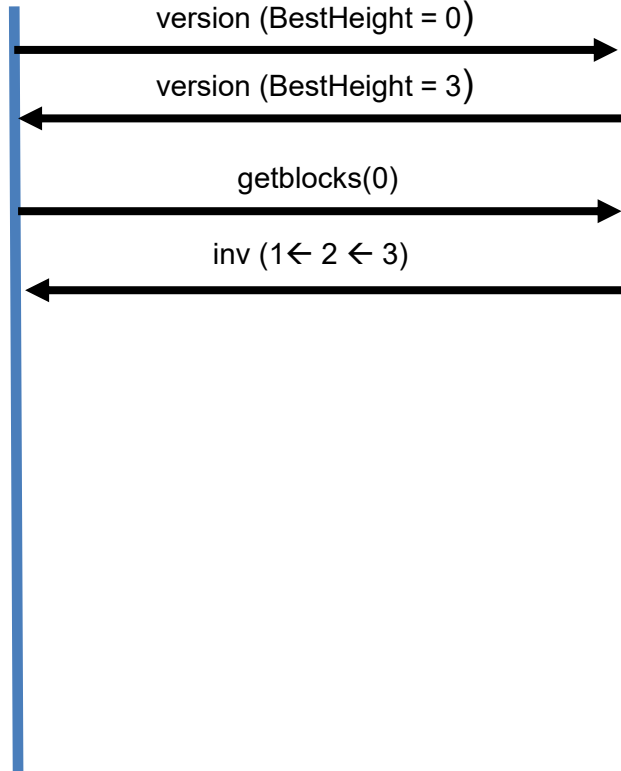


# Initial block download

0

node A

node B



0

1

2

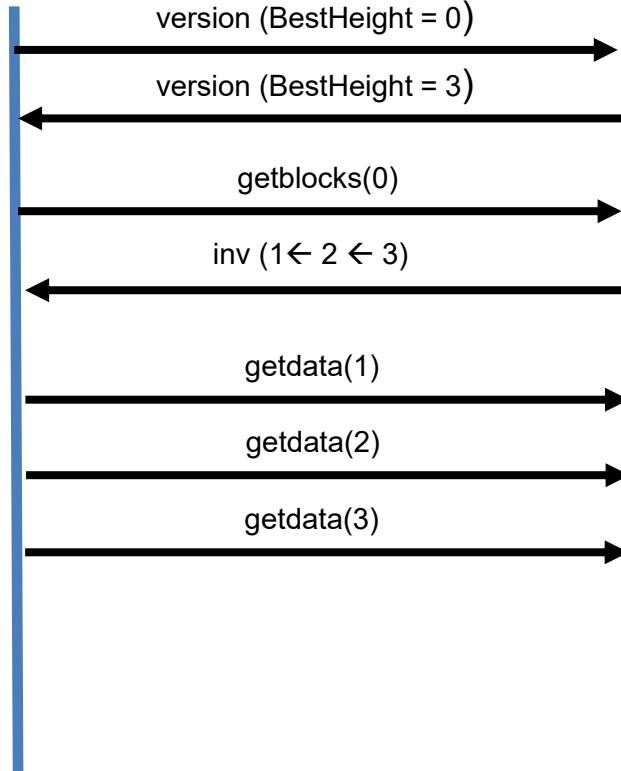
3

# Initial block download

0

node A

node B



0

1

2

3

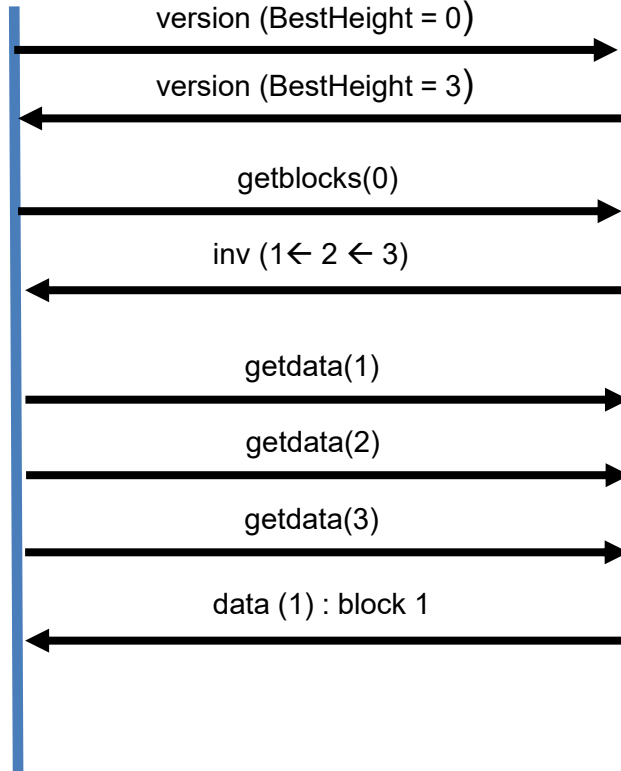


# Initial block download

0

node A

node B

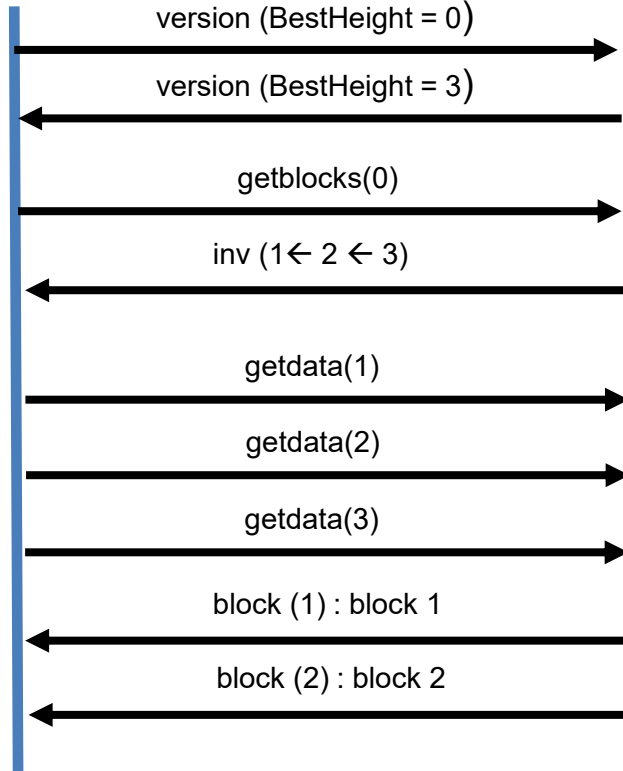


# Initial block download

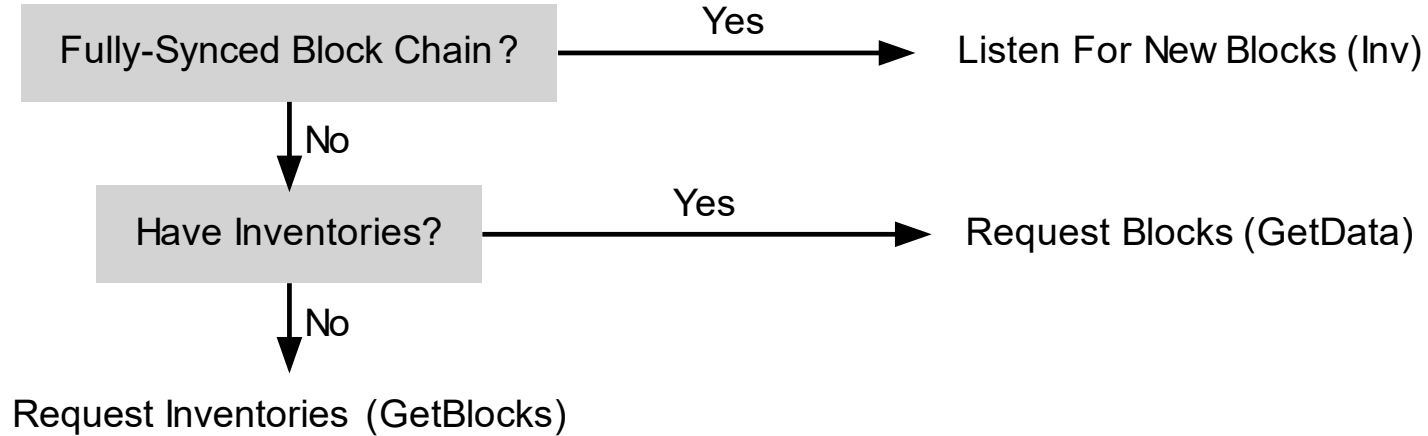
0

node A

node B

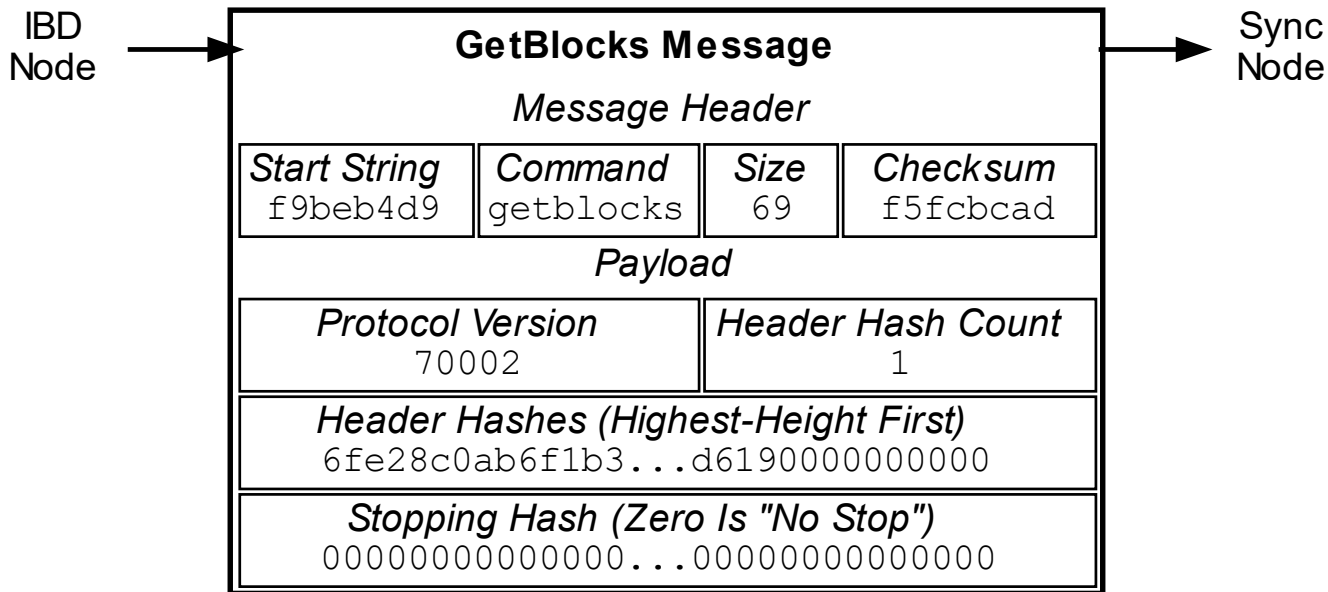


# In reality (IBD)



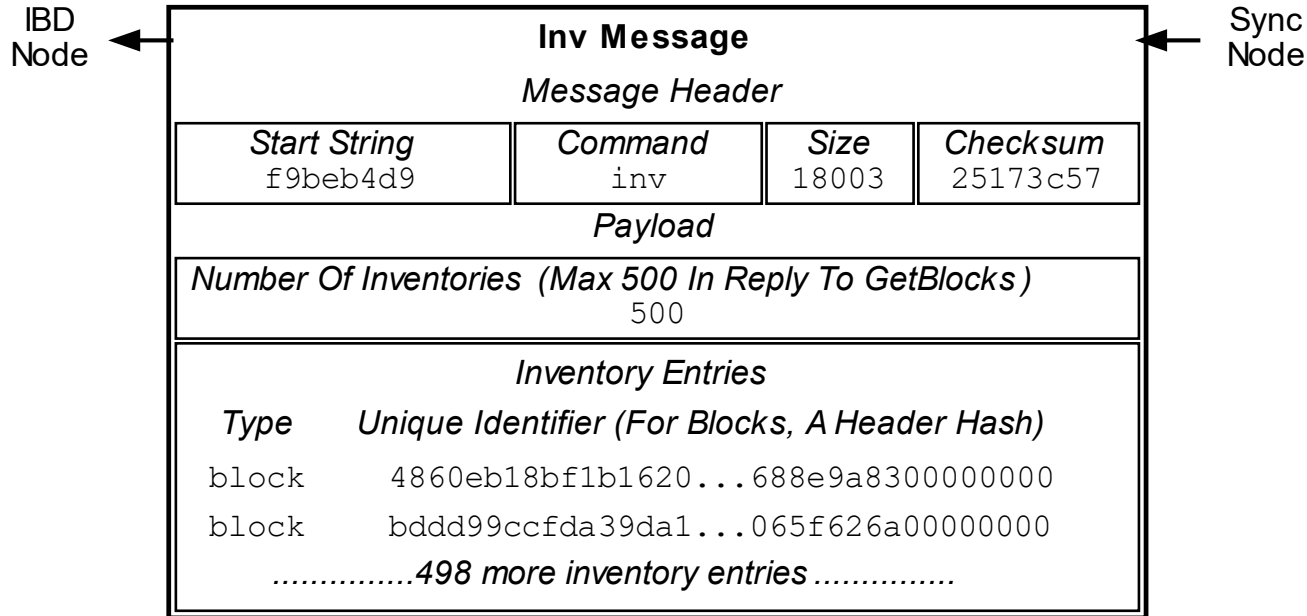
Overview Of Blocks-First Initial Blocks Download (IBD)

# In reality (IBD)



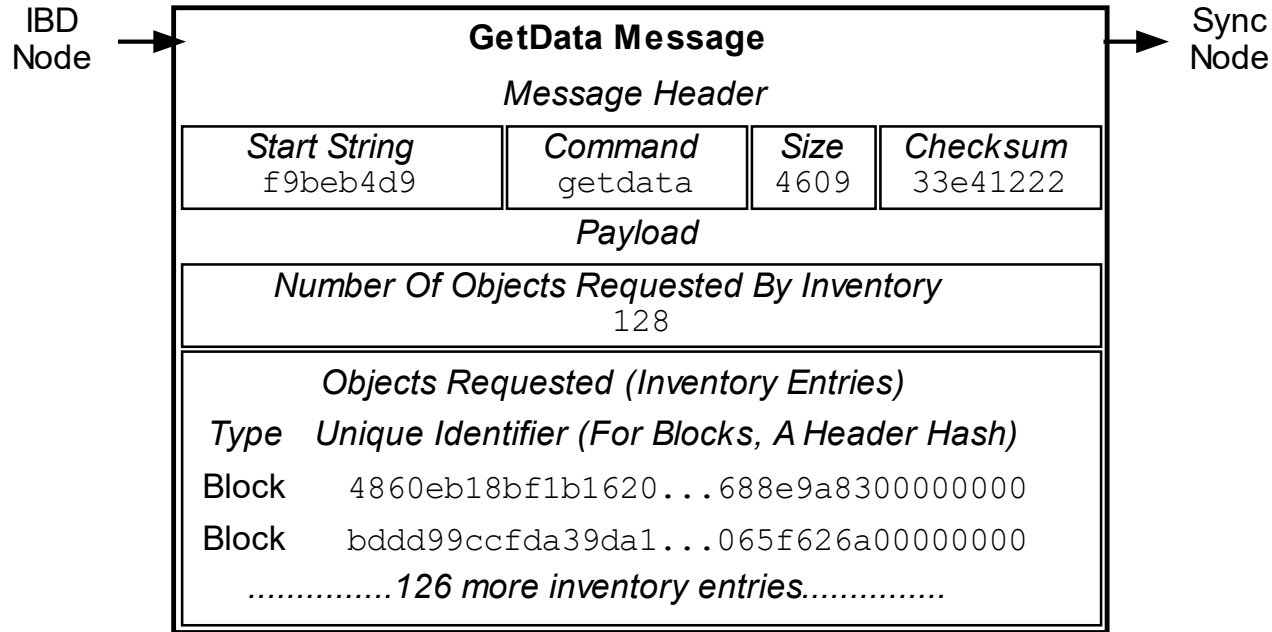
First getblocks message sent from Initial Blocks Download (IBD) node

# In reality (IBD)



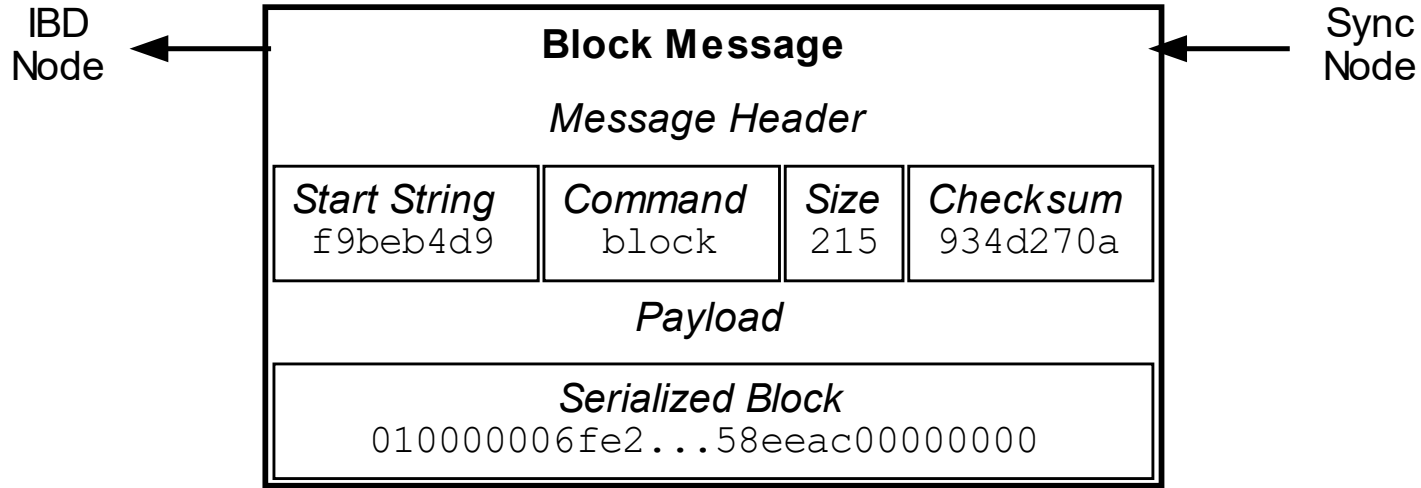
First inv message reply sent to Initial Blocks Download (IBD) node

# In reality (IBD)



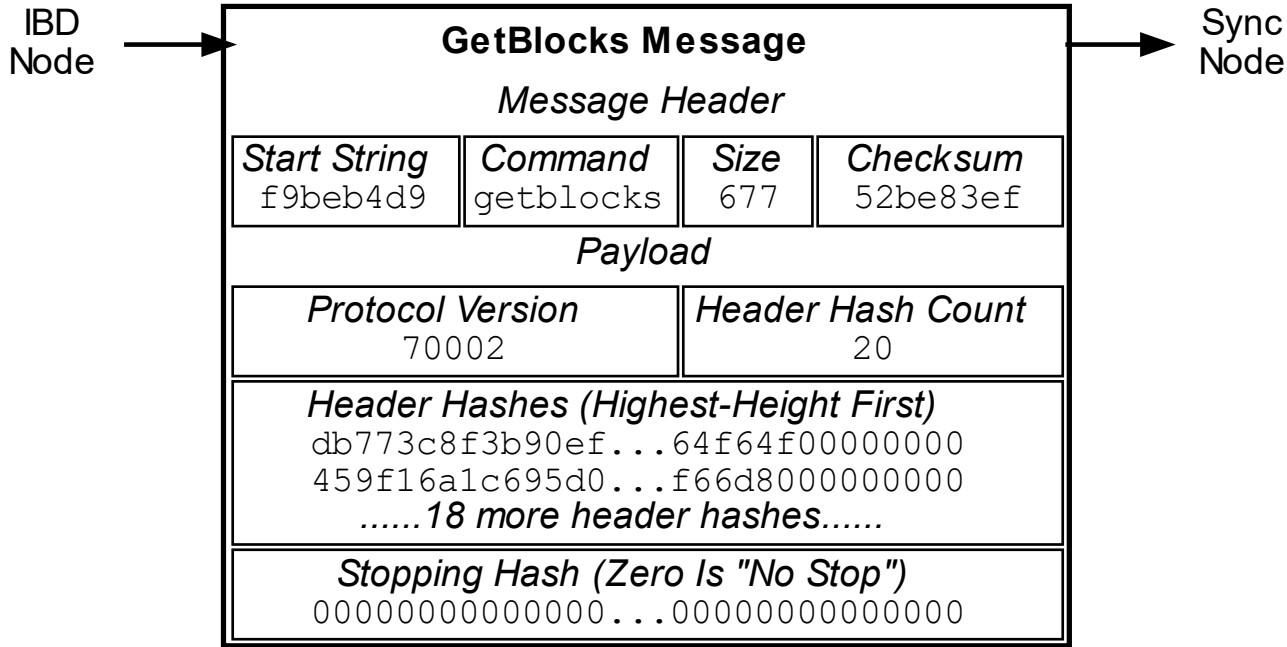
First getdata message sent from Initial Blocks Download (IBD) node

# In reality (IBD)



First block message sent to Initial Blocks Download (IBD) node

# In reality (IBD)



Second getblocks message sent from Initial Blocks Download (IBD) node



## **Simple payment verification (SPV):**

- A typical user is only interested in transactions using her own address
- A typical use of a Bitcoin wallet
- SPV nodes do not store the entire blockchain
- They store only the block headers (80 bytes per block = 4.2MB per year)
- Much smaller than a full node
- But they can not validate all the UTXOs

**An SPV node can verify only two things:**

1. That a transaction *tx* belongs to the block *blockX*
2. That a block *blockX* has k confirmations

## How does an SPV node verify that tx belongs to blockX?

- An SPV node store the block headers

Field	Purpose	Updated when...	Size (Bytes)
Version	Block version number	You upgrade the software and it specifies a new version	4
hashPrevBlock	256-bit hash of the previous block header	A new block comes in	32
hashMerkleRoot	256-bit hash based on all of the transactions in the block	A transaction is accepted	32
Time	Current timestamp as seconds since 1970-01-01T00:00 UTC	Every few seconds	4
Bits	Current <a href="#">target</a> in compact format	The <a href="#">difficulty</a> is adjusted	4
Nonce	32-bit number (starts at 0)	A hash is tried (increments)	4

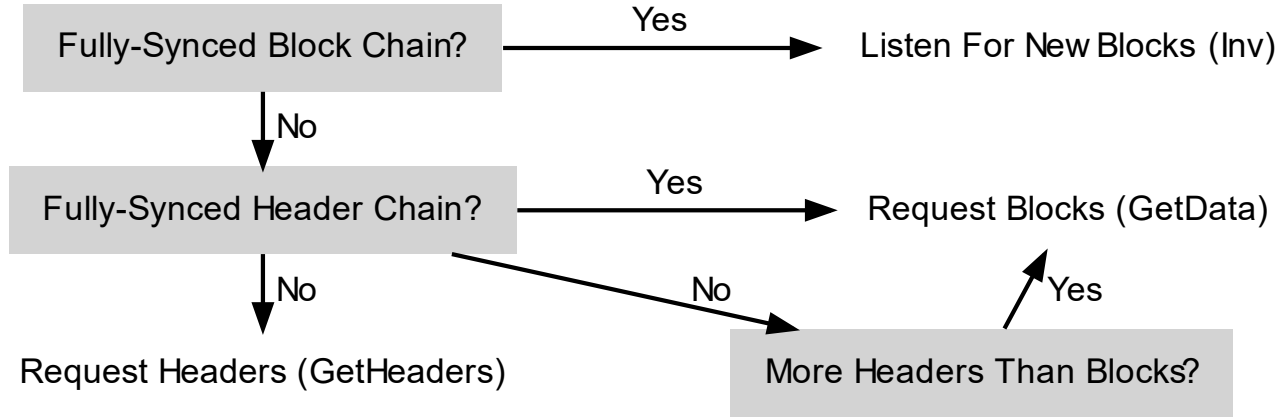
## How does an SPV node verify that tx belongs to blockX?

- An SPV node store the block headers
- An SPV node A asks a full node B for a certificate for *tx*
- And validates the certificate against *hashMerkleRoot*

**How does an SPV node A verify that tx has k confirmations?**

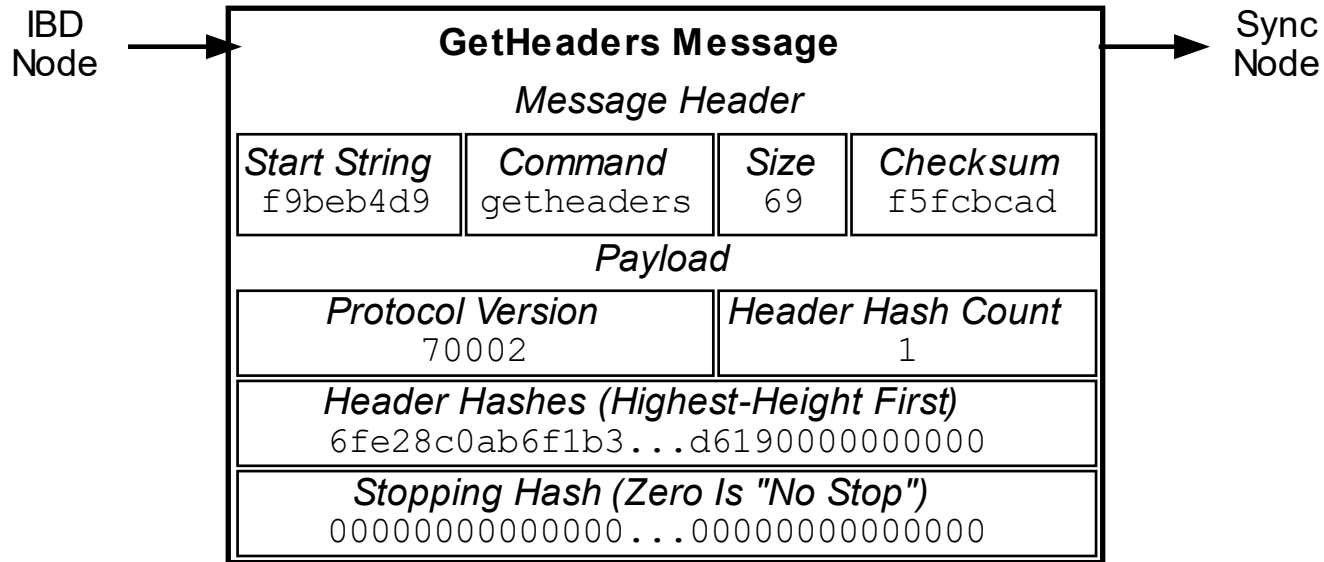
- A has the headers up until the current end of blockchain
- Can A verify that the headers are valid?

# IBD for an SPV node



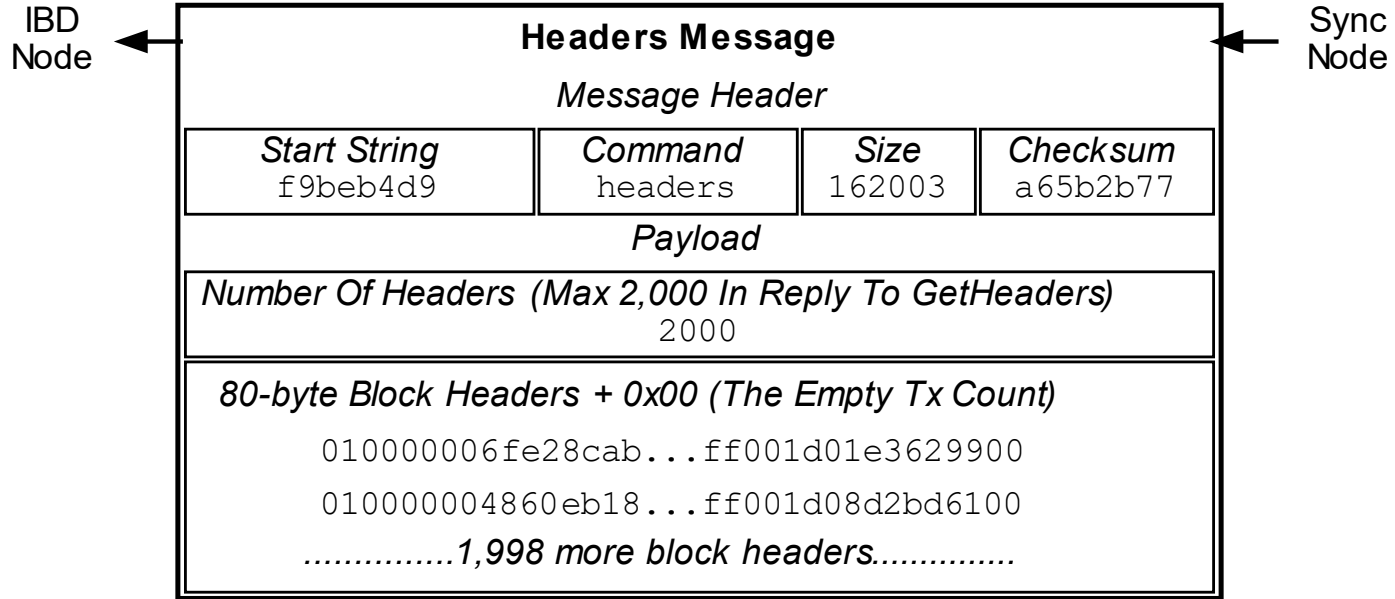
Overview Of Headers-First Initial Blocks Download (IBD)

# IBD for an SPV node



First getheaders message sent from Initial Blocks Download (IBD) node

# IBD for an SPV node



First headers message reply sent to Initial Blocks Download (IBD) node



# Vulnerabilities of SPV nodes

## First vulnerability: partial DOS attack

- If a full node is malicious can it pretend that *tx* belongs to *blockX*?
- Can it pretend that *tx* does not belong to *blockX*?

## Solution:

- The SPV node should use more than one full node
- Merkle trees should be ordered

# Vulnerabilities of SPV nodes

## **Second vulnerability: loss of privacy**


- If the SPV node queries node B a lot
- B will know all the addresses of A
- It can track A
- It can prohibit A to broadcast transactions
- How to solve this?

# Bloom filters

## **A probabilistic search filter:**

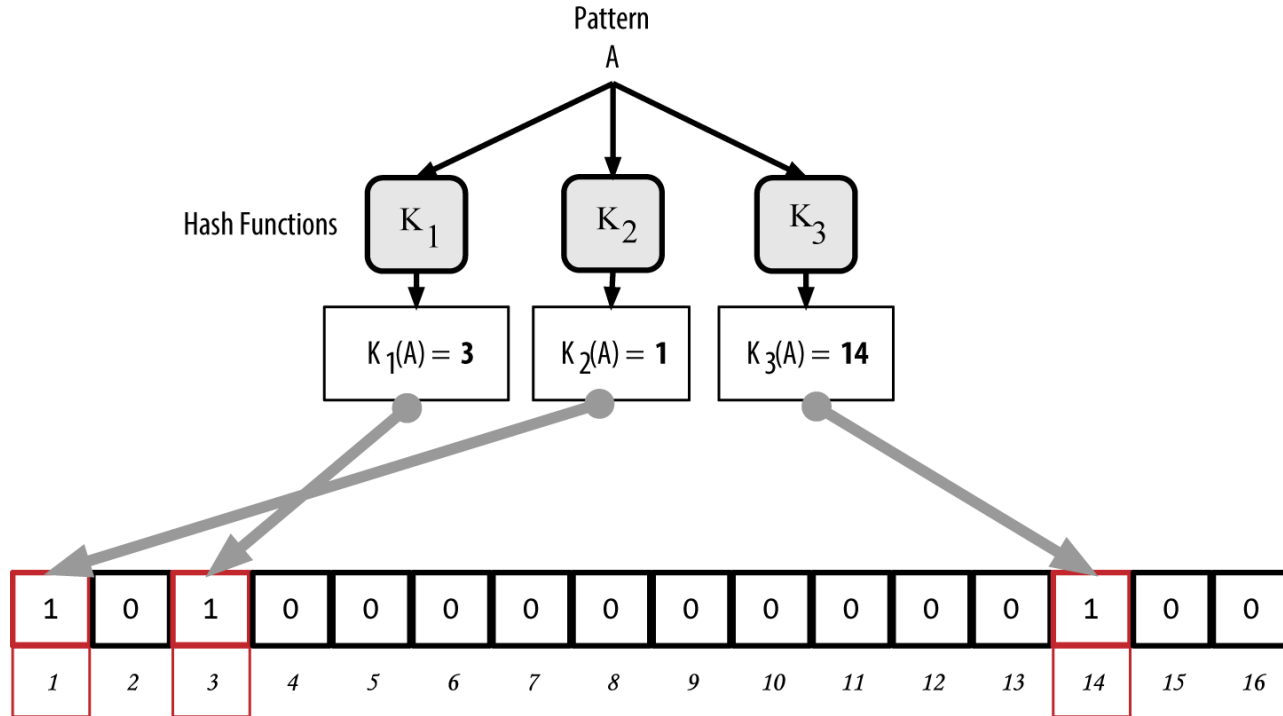
- The filter starts empty
- Elements are added to the filter
- One wishes to check if an element was added to the filter
- Can give false positives
- Can not give false negatives

# Bloom filters



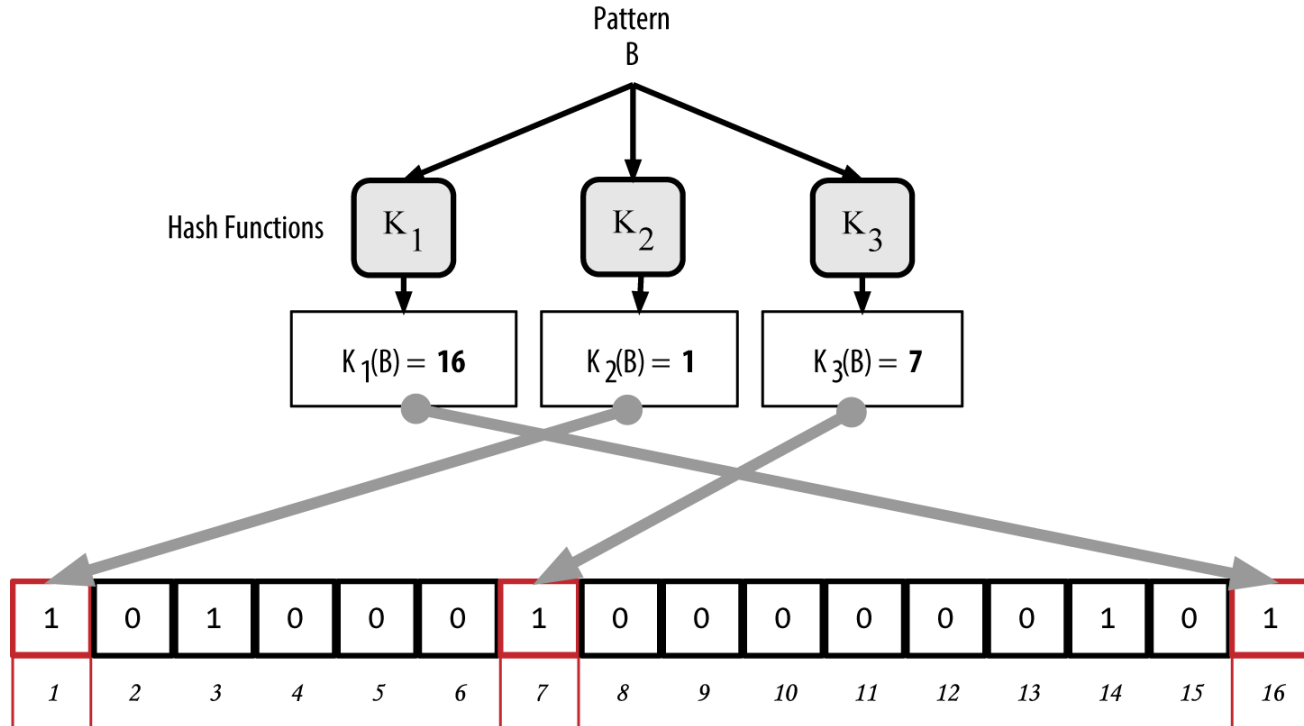
# Bloom filters

Inserción



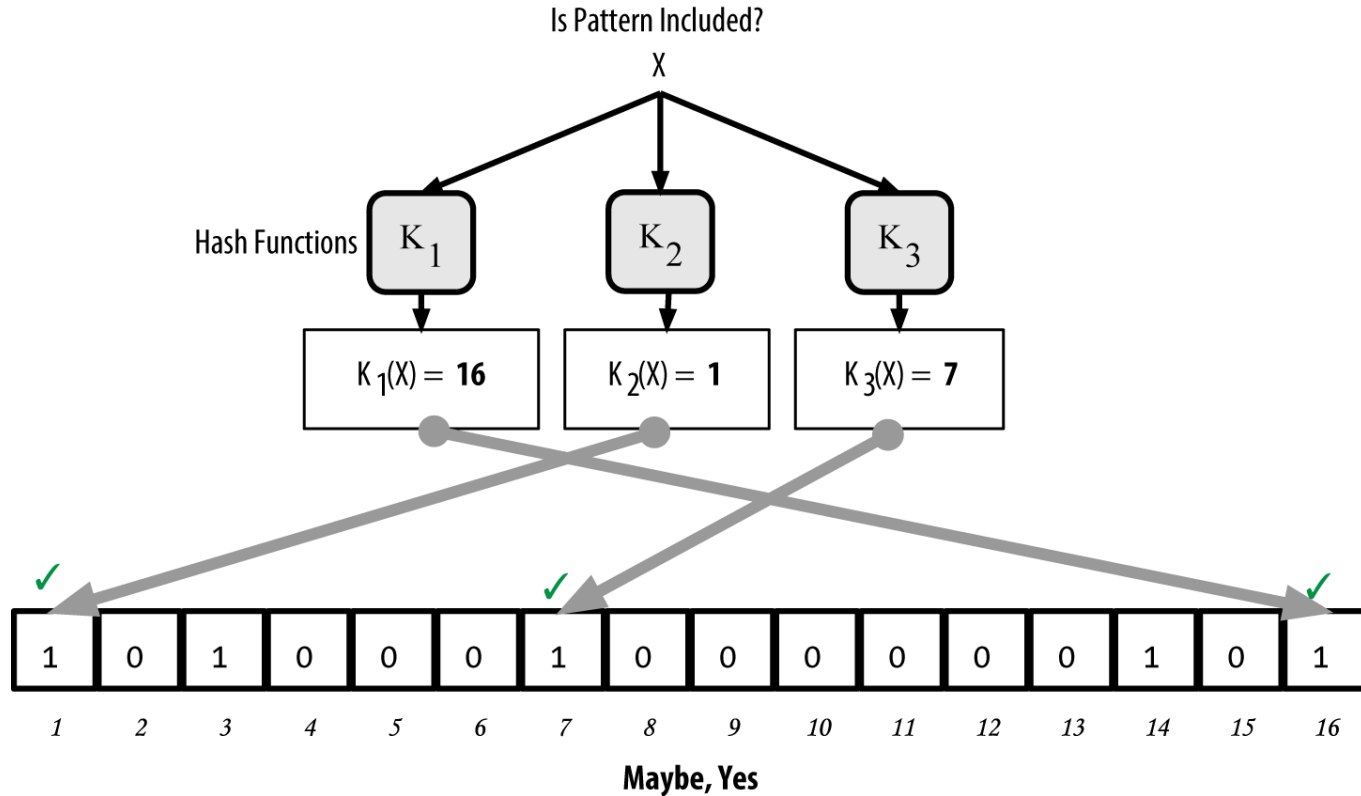
# Bloom filters

Insertion



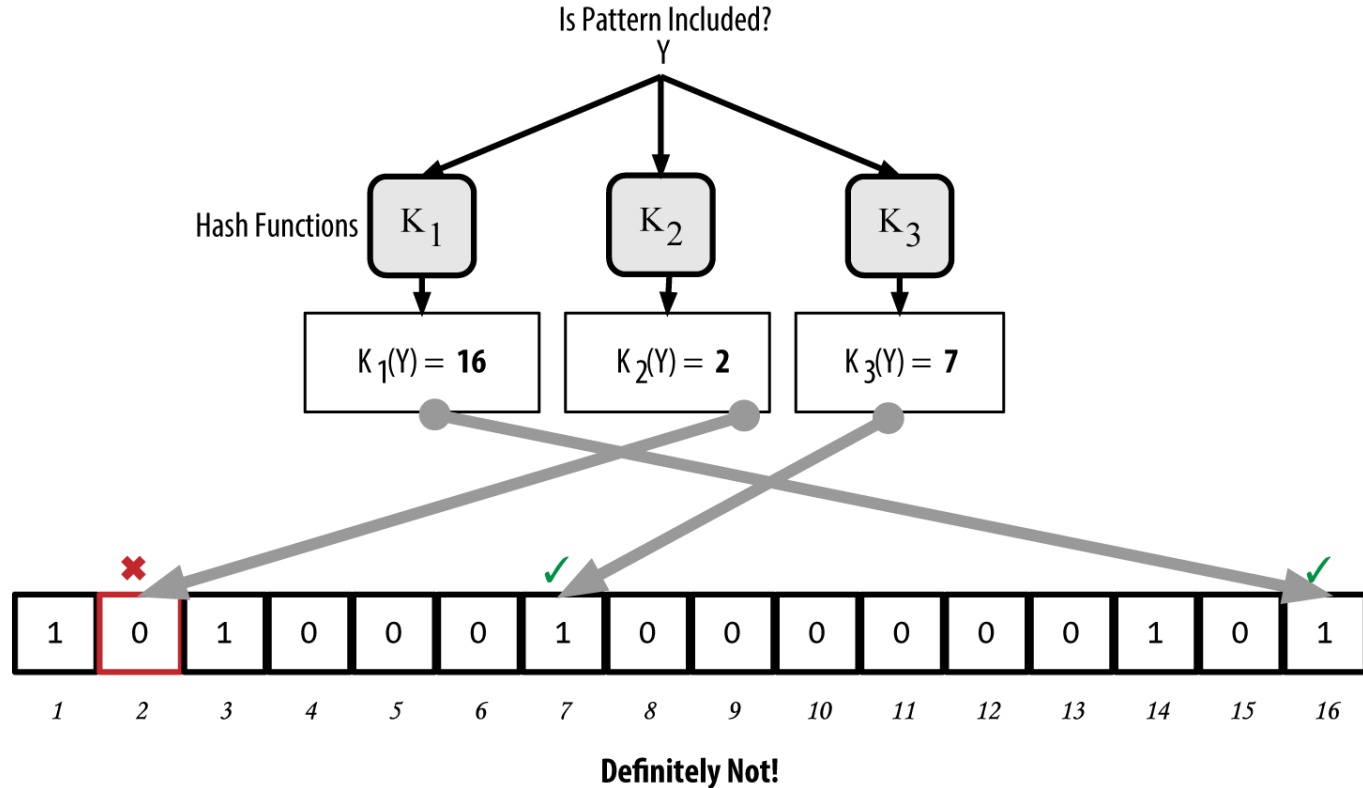
# Bloom filters

Search



# Bloom filters

Search





# SPV nodes and Bloom filters

## SPV A and its full node B:

- A inserts the UTXOs and addresses of interest to a Bloom filter (BF)
- A inserts a bit of junk data into the BF
- A sends its BF to B
- B uses the BF in the communication channel with A
- For each *tx* that B receives, if *tx* matches the BF, *tx* is sent to A
- A validates whether *tx* was a false positive and in this case it discards *tx*

# SPV nodes and Bloom filters

## What B checks in a transaction:

- Hash of the transaction (the ID)
- Inputs
- Outputs
- All the addresses

## Read more on this:

- <https://github.com/bitcoin/bips/blob/master/bip-0037.mediawiki#filter-matching-algorithm>

# SPV nodes and Bloom filters

## Numbers used in Bitcoin:

- The size of the filter is 36.000 bytes max
- 50 hash functions
- Full node sends the header + the certificate for the transactions matching the filter

## Read more on this:

- <https://github.com/bitcoin/bips/blob/master/bip-0037.mediawiki#filter-matching-algorithm>

# Bitcoin network

## References:

1. Antonopoulos: <https://github.com/bitcoinbook/bitcoinbook/blob/develop/ch08.asciidoc>
2. Protocol specification: [https://en.bitcoin.it/wiki/Protocol\\_documentation](https://en.bitcoin.it/wiki/Protocol_documentation)
3. Network: <https://en.bitcoin.it/wiki/Network>
4. P2P network: <https://bitcoin.org/en/developer-guide#peer-discovery>
5. SPV nodes: <https://bitcoin.org/en/developer-guide#simplified-payment-verification-spv>
6. Bloom filters: <https://github.com/bitcoin/bips/blob/master/bip-0037.mediawiki#filter-matching-algorithm>

# Mining pools

## **What does a miner do on her own:**

- Buys hardware
- Spends electricity to compute the hashes
- Gets Bitcoins (ocassionally)
- A huge variance on when it will discover the blocks
- It can happen it does not discover a block for two years
- Not a sustainable bussiness strategy

# Mining pools

## Solution:

- Several miners join a *mining pool*
  - Similar to a producers collective (i.e. zadruga)
  - They all solve the same mining puzzle
  - *Pool manager* gives them the puzzle = block
- 
- How does a pool manager know how much to pay each miner?
  - Proportional to the work the miner did
  - How can we measure this?

# Mining pools

## Mining shares:

- An incomplete solution to the mining puzzle
- Puzzle to mine the block has the difficulty of 70 zeros
- A share has a difficulty of 50 zeros
  
- Each time a miner mines a share, it sends it to the pool manager
- Manager pays in proportion with the number of mined shares
  
- It can happen that the miner that mined the block receives less than others

# Mining pools

## Mining shares:

- An incomplete solution to the mining puzzle
- Puzzle to mine the block has the difficulty of 70 zeros
- A share has a difficulty of 50 zeros

```
4AA087F0A52ED2093FA816E53B9B6317F9B8C1227A61F9481AFED67301F2E3FB
D3E51477DCAB108750A5BC9093F6510759CC880BB171A5B77FB4A34ACA27DEDD
00000000008534FF68B98935D090DF5669E3403BD16F1CD4D41CF17D6B474255
BB34ECA3DBB52EFF4B104EBBC0974841EF2F3A59EBBC4474A12F9F595EB81F4B
00000000002F891C1E232F687E41515637F7699EA0F462C2564233FE082BB0AF
0090488133779E7E98177AF1C765CF02D01AB4848DF555533B6C4CFCA201CBA1
460BEFA43B7083E502D36D9D08D64AFB99A100B3B80D4EA4F7B38E18174A0BFB
000000000000000078FB7E1F7E2E4854B8BC71412197EB1448911FA77BAE808A
652F374601D149AC47E01E7776138456181FA4F9D0EEDD8C4FDE3BEF6B1B7ECE
785526402143A291CFD60DA09CC80DD066BC723FD5FD20F9B50D614313529AF3
000000000041EE593434686000AF77F54CDE839A6CE30957B14EDEC10B15C9E5
9C20B06B01A0136F192BD48E0F372A4B9E6BA6ABC36F02FCED22FD9780026A8F
```




# Mining pools

## Mining shares:

- An incomplete solution to the mining puzzle
- Puzzle to mine the block has the difficulty of 70 zeros
- A share has a difficulty of 50 zeros
  
- Each time a miner mines a share, it sends it to the pool manager
- Manager pays in proportion with the number of mined shares
  
- It can happen that the miner that mined the block receives less than others

# Mining pools

It can happen that the miner that mined the block receives less than others



# Mining pools

## Operating a pool:

- How to pay the miners?
- Op1: flat fee – manager assumes the risk (pays even no block is mined)
- Op2: proportional – manager pays when a block is mined
- Miners can play with all these options

# Mining pools

## **Mining pools and the 51% attack**


- Ghash.io – in 2014 had 50% of the hash power
- The market autoregulates (at least that's what capitalism is based on)

## **Two schools of thought:**

1. Bitcoin mining is wasteful
2. Bitcoin mining is environmentally friendly in the long term


# Network recap

Node functionality



# Network recap

## Node type



### Reference Client (Bitcoin Core)

Contains a Wallet, Miner, full Blockchain database, and Network routing node on the bitcoin P2P network.




### Full Block Chain Node

Contains a full Blockchain database, and Network routing node on the bitcoin P2P network.



### Solo Miner

Contains a mining function with a full copy of the blockchain and a bitcoin P2P network routing node.



### Lightweight (SPV) wallet

Contains a Wallet and a Network node on the bitcoin P2P protocol, without a blockchain.




### Pool Protocol Servers

Gateway routers connecting the bitcoin P2P network to nodes running other protocols such as pool mining nodes or Stratum nodes.



### Mining Nodes

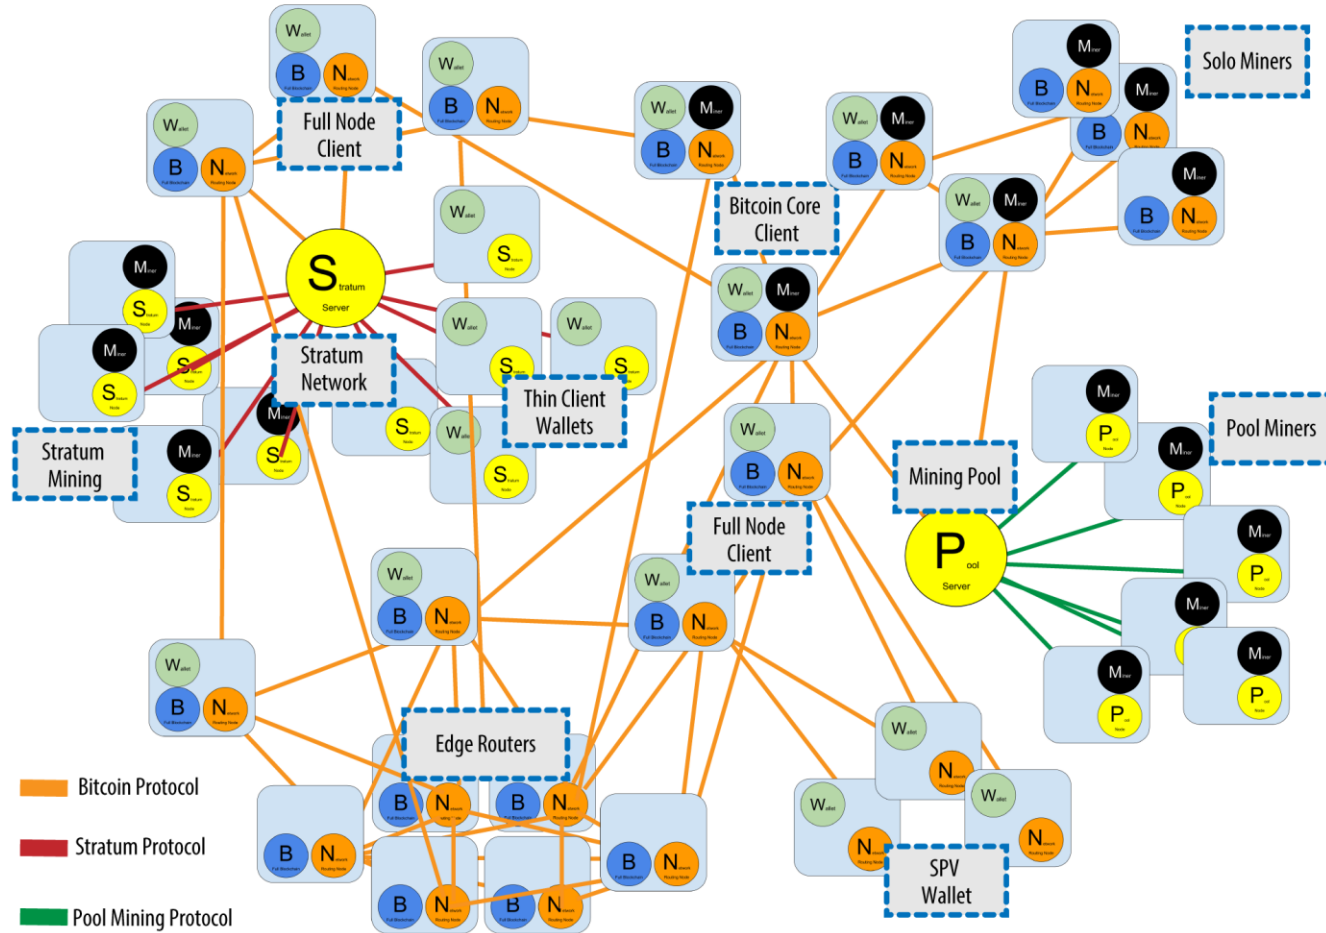
Contain a mining function, without a blockchain, with the Stratum protocol node (S) or other pool (P) mining protocol node.



### Lightweight (SPV) Stratum wallet

Contains a Wallet and a Network node on the Stratum protocol, without a blockchain.

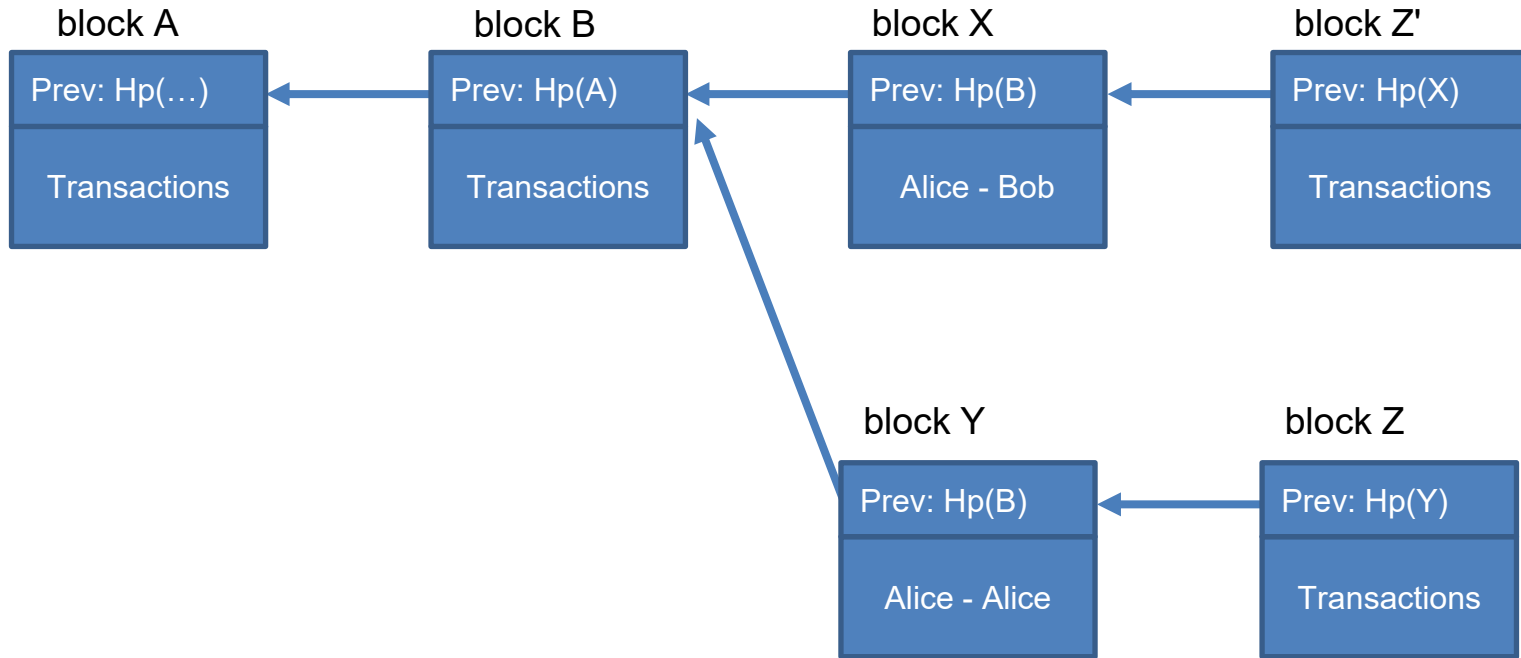
# Network recap





# Forks

**Two (or more) branches in the blockchain:**



## **A rule change**

- All the nodes need to upgrade their software
- Does not happen naturally
- **Hard forks vs. Soft forks**

## Hard fork:

- A change that is not forward compatible
- Introduces rules that were not valid previously
- E.g. you should sign double hash of a transaction  $H(H(tx))$ , and not only  $H(tx)$
- Nodes that did the upgrade are OK, but the rest are not
- Blockchain divides
- Example: Bitcoin Cash

## **Soft fork:**

- A change that is forward compatible
- Introduces stricter rules
- Nodes with the old software will accept all new blocks
- Nodes with the new software will reject some blocks that were previously valid
- Miners with the old software can realize that they need to upgrade
- Example: (almost) any BIP; e.g. P2SH